



Computer Vision; Image Transformation; Optical Flow and Depth Estimation

Maziar Raissi

Assistant Professor

Department of Applied Mathematics

University of Colorado Boulder

maziar.raissi@colorado.edu



Boulder

FlowNet: Learning Optical Flow with Convolutional Networks

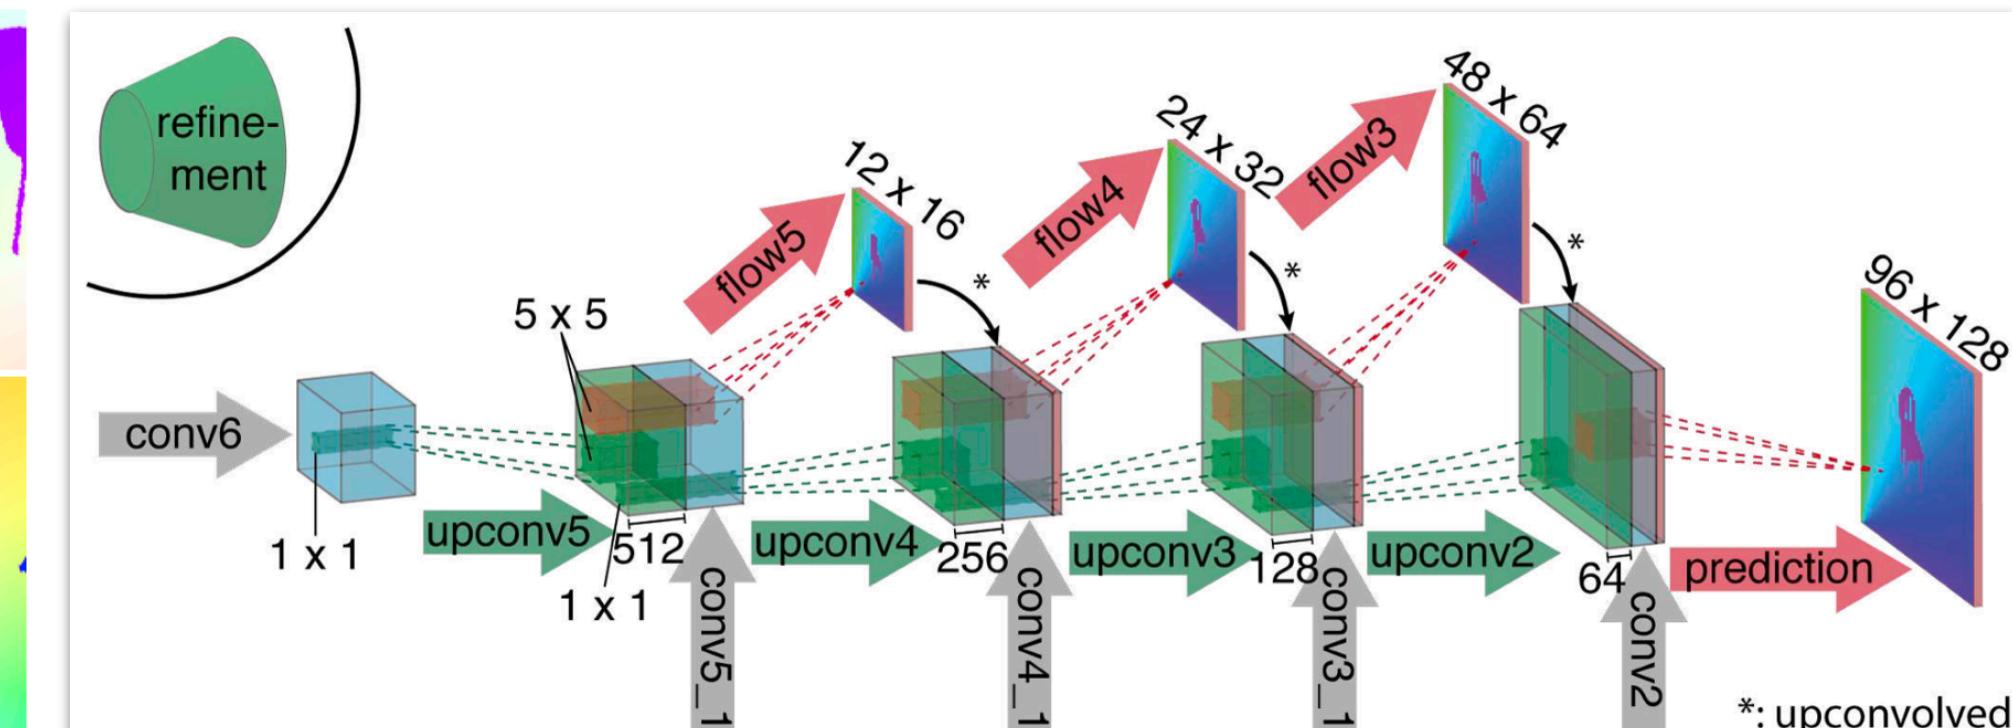
Given a dataset consisting of image pairs and ground truth flows, train a network to predict the $x-y$ flow fields directly from the images.

Flying Chairs Dataset

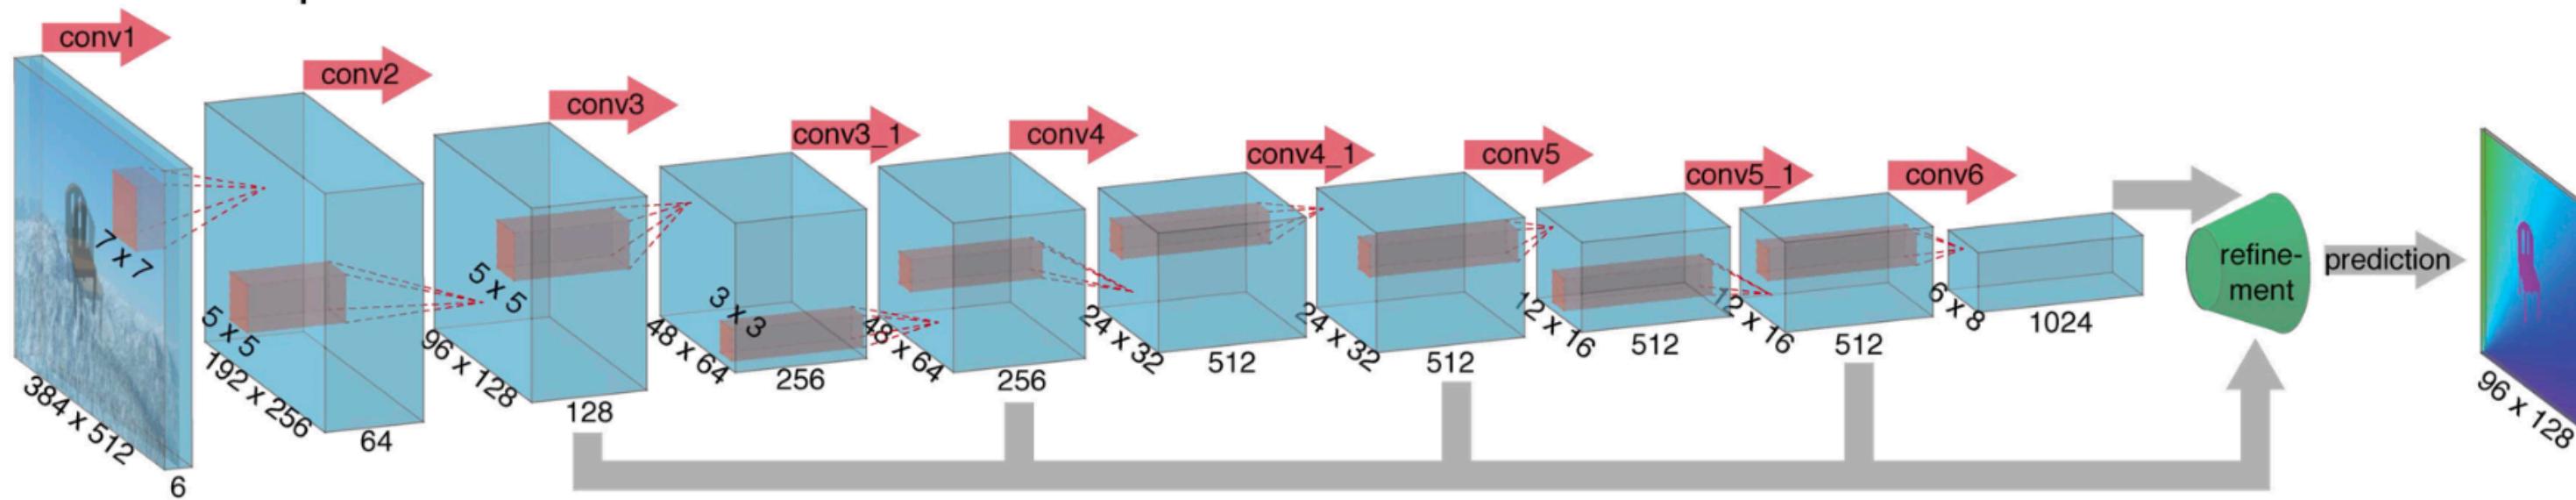


Data Augmentation

upconv: unpooling + conv



FlowNetSimple



Method	Sintel Clean		Sintel Final		KITTI		Middlebury train		Middlebury test		Chairs	Time (sec)
	train	test	train	test	train	test	AEE	AAE	AEE	AAE		
EpicFlow [30]	2.27	4.12	3.57	6.29	3.47	3.8	0.31	3.24	0.39	3.55	2.94	16 -
DeepFlow [35]	3.19	5.38	4.40	7.21	4.58	5.8	0.21	3.04	0.42	4.22	3.53	17 -
EPPM [3]	-	6.49	-	8.38	-	9.2	-	-	0.33	3.36	-	0.2
LDOF [6]	4.19	7.56	6.28	9.12	13.73	12.4	0.45	4.97	0.56	4.55	3.47	65 2.5
FlowNetS	4.50	7.42	5.45	8.43	8.26	-	1.09	13.28	-	-	2.71	- 0.08
FlowNetS+v	3.66	6.45	4.76	7.67	6.50	-	0.33	3.87	-	-	2.86	- 1.05
FlowNetS+ft	(3.66)	6.96	(4.44)	7.76	7.52	9.1	0.98	15.20	-	-	3.04	- 0.08
FlowNetS+ft+v	(2.97)	6.16	(4.07)	7.22	6.07	7.6	0.32	3.84	0.47	4.58	3.03	- 1.05

	Frame pairs	Frames with ground truth	Ground truth density per frame
Middlebury	72	8	100%
KITTI	194	194	~50%
Sintel	1,041	1,041	100%
Flying Chairs	22,872	22,872	100%



End Point Error (EPE) Loss: Euclidean distance between the predicted flow vector and the ground truth, averaged over all pixels



Boulder

FlowNet 2.0: Evolution of Optical Flow Estimation with Deep Networks



- problems with small displacements
- noisy artifacts in estimated flow fields

Dataset Schedules

Numbers indicate endpoint errors on Sintel train clean

Architecture	Datasets	S_{short}	S_{long}	S_{fine}
FlowNetS	Chairs	4.45	-	-
	Chairs	-	4.24	4.21
	Things3D	-	5.07	4.50
	mixed	-	4.52	4.10
	Chairs → Things3D	-	4.24	3.79
FlowNetC	Chairs	3.77	-	-
	Chairs → Things3D	-	3.58	3.04

Training on Chairs first and fine-tuning on Things3D yields the best results. FlowNetC performs better than FlowNetS.

Stacking Networks

$I_1, I_2 \rightarrow$ images $w = (u, v)^T \rightarrow$ previous flow estimation

$\tilde{I}_2(x, y) = I_2(x + u, y + v) \rightarrow$ warp the second image I_2 (bilinear interpolation)

$e = \|\tilde{I}_2(x, y) - I_1\| \rightarrow$ error (\tilde{I}_2 should be close to I_1)

FlowNetCorr

$$f_1, f_2 \in \mathbb{R}^{w \times h \times c} \text{ or } f_1, f_2 : \mathbb{R}^2 \rightarrow \mathbb{R}^c$$

$c(x_1, x_2) \rightarrow$ correlation of two patches centered at x_1 in the first map and x_2 in the second map

$$c(x_1, x_2) = \sum_{o \in [-k, k] \times [-k, k]} \langle f_1(x_1 + o), f_2(x_2 + o) \rangle \rightarrow \text{for a square patch of size } K = 2k + 1$$

$$w^2 h^2 c K^2 \rightarrow \text{computational cost}$$

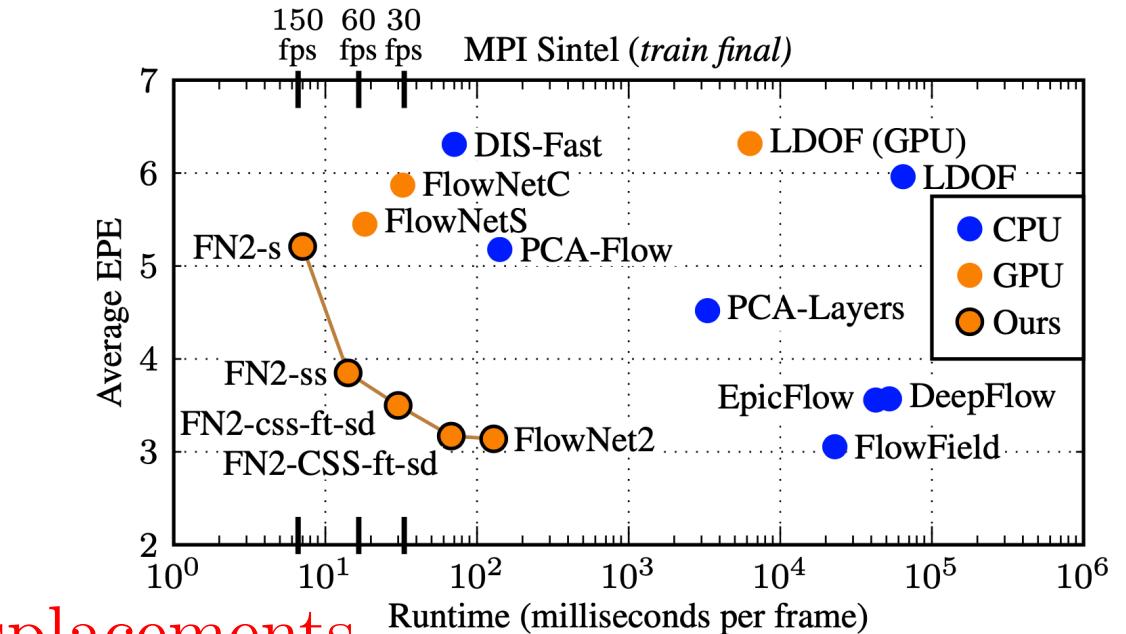
compute $c(x_1, x_2)$ only in a neighborhood of x_1 of size $D = 2d + 1$

$$w h D^2 c K^2 \rightarrow \text{computational cost}$$

$$w \times h \times D^2 \rightarrow \text{output size}$$

use stride s_1 to quantize x_1 globally

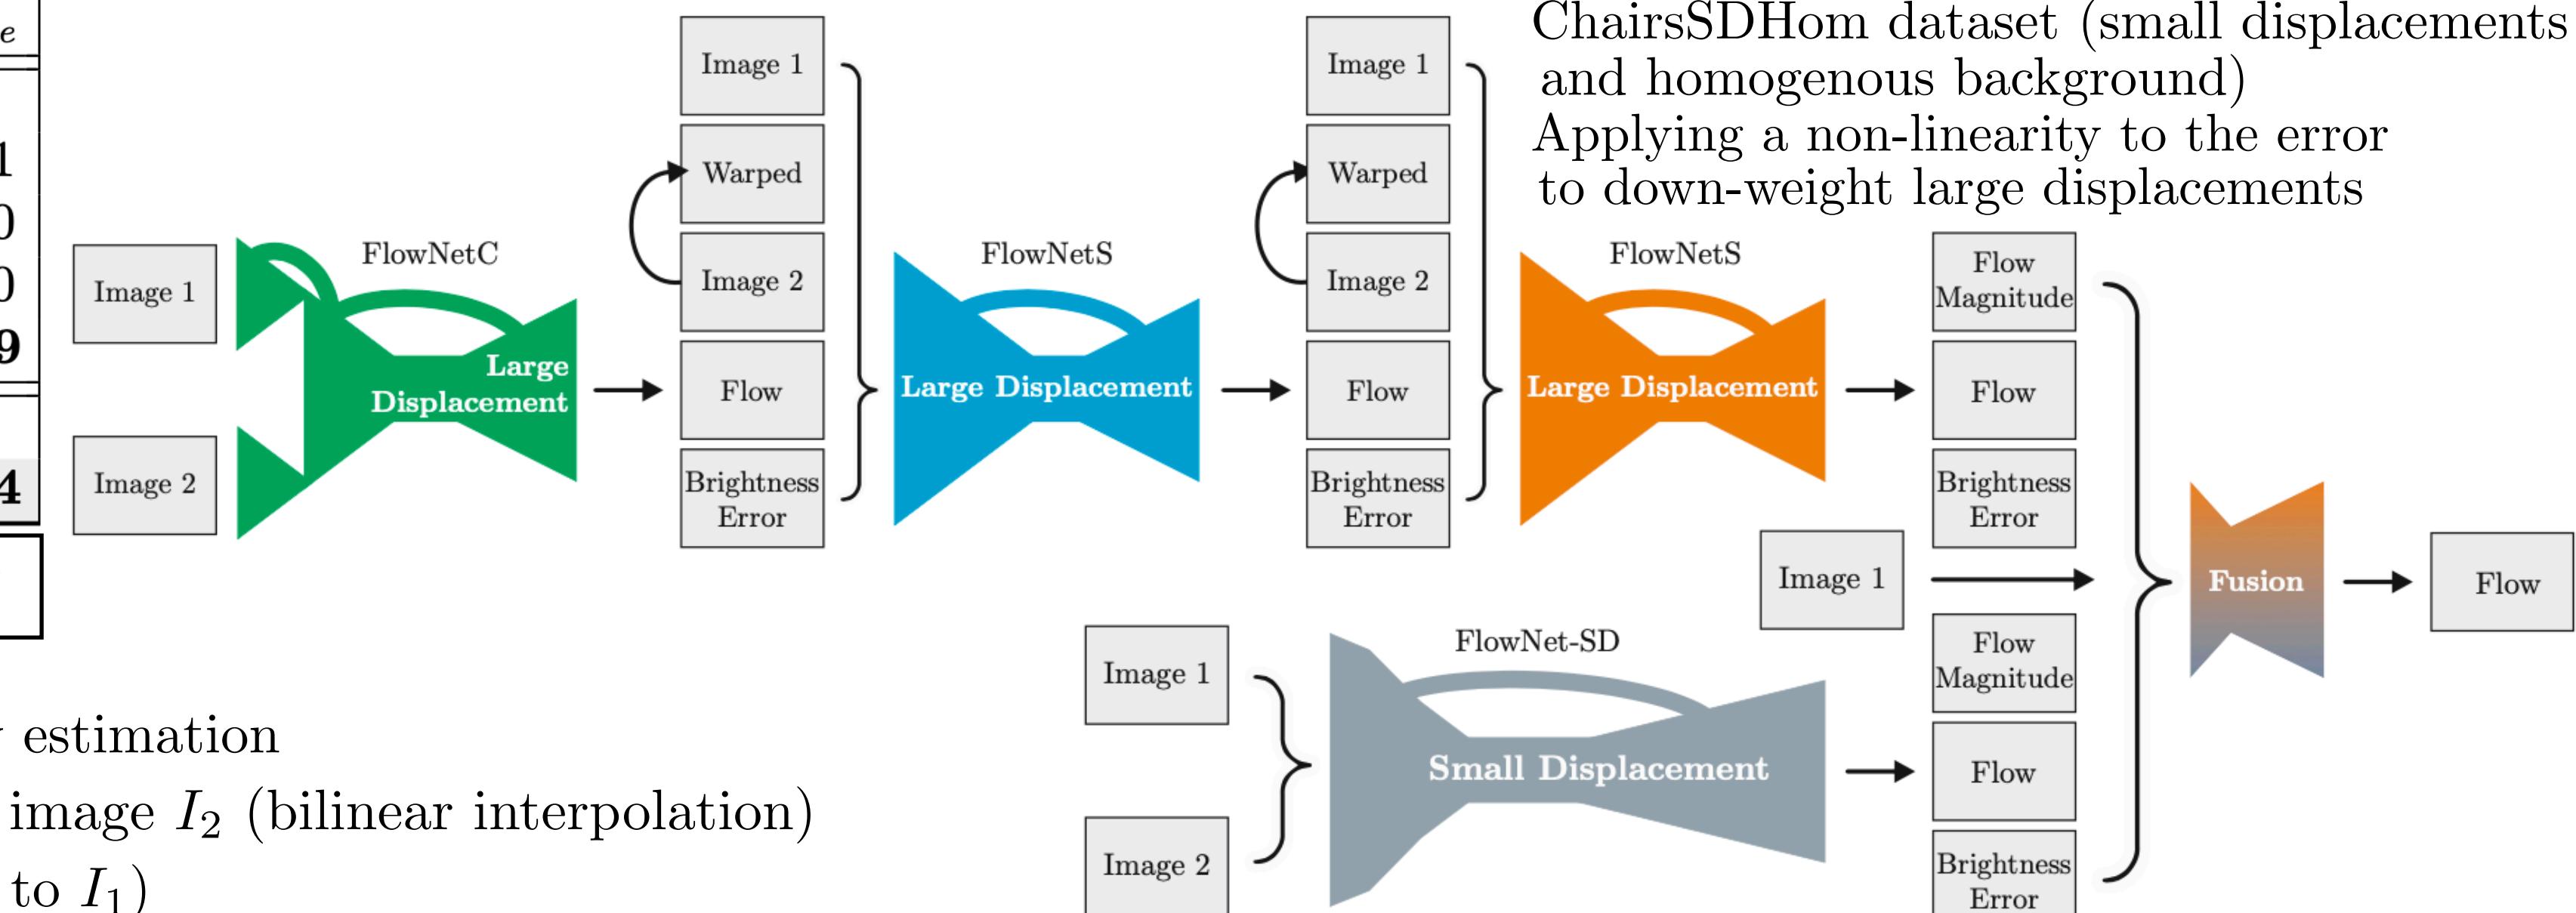
use stride s_2 to quantize x_2 locally around x_1

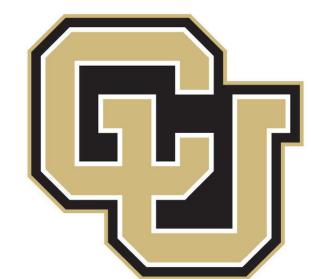


Small Displacements

ChairsSDHom dataset (small displacements and homogenous background)

Applying a non-linearity to the error to down-weight large displacements





Boulder

Questions?
