



Boulder

Computer Vision; 3D



[YouTube Playlist](#)

Maziar Raissi

Assistant Professor

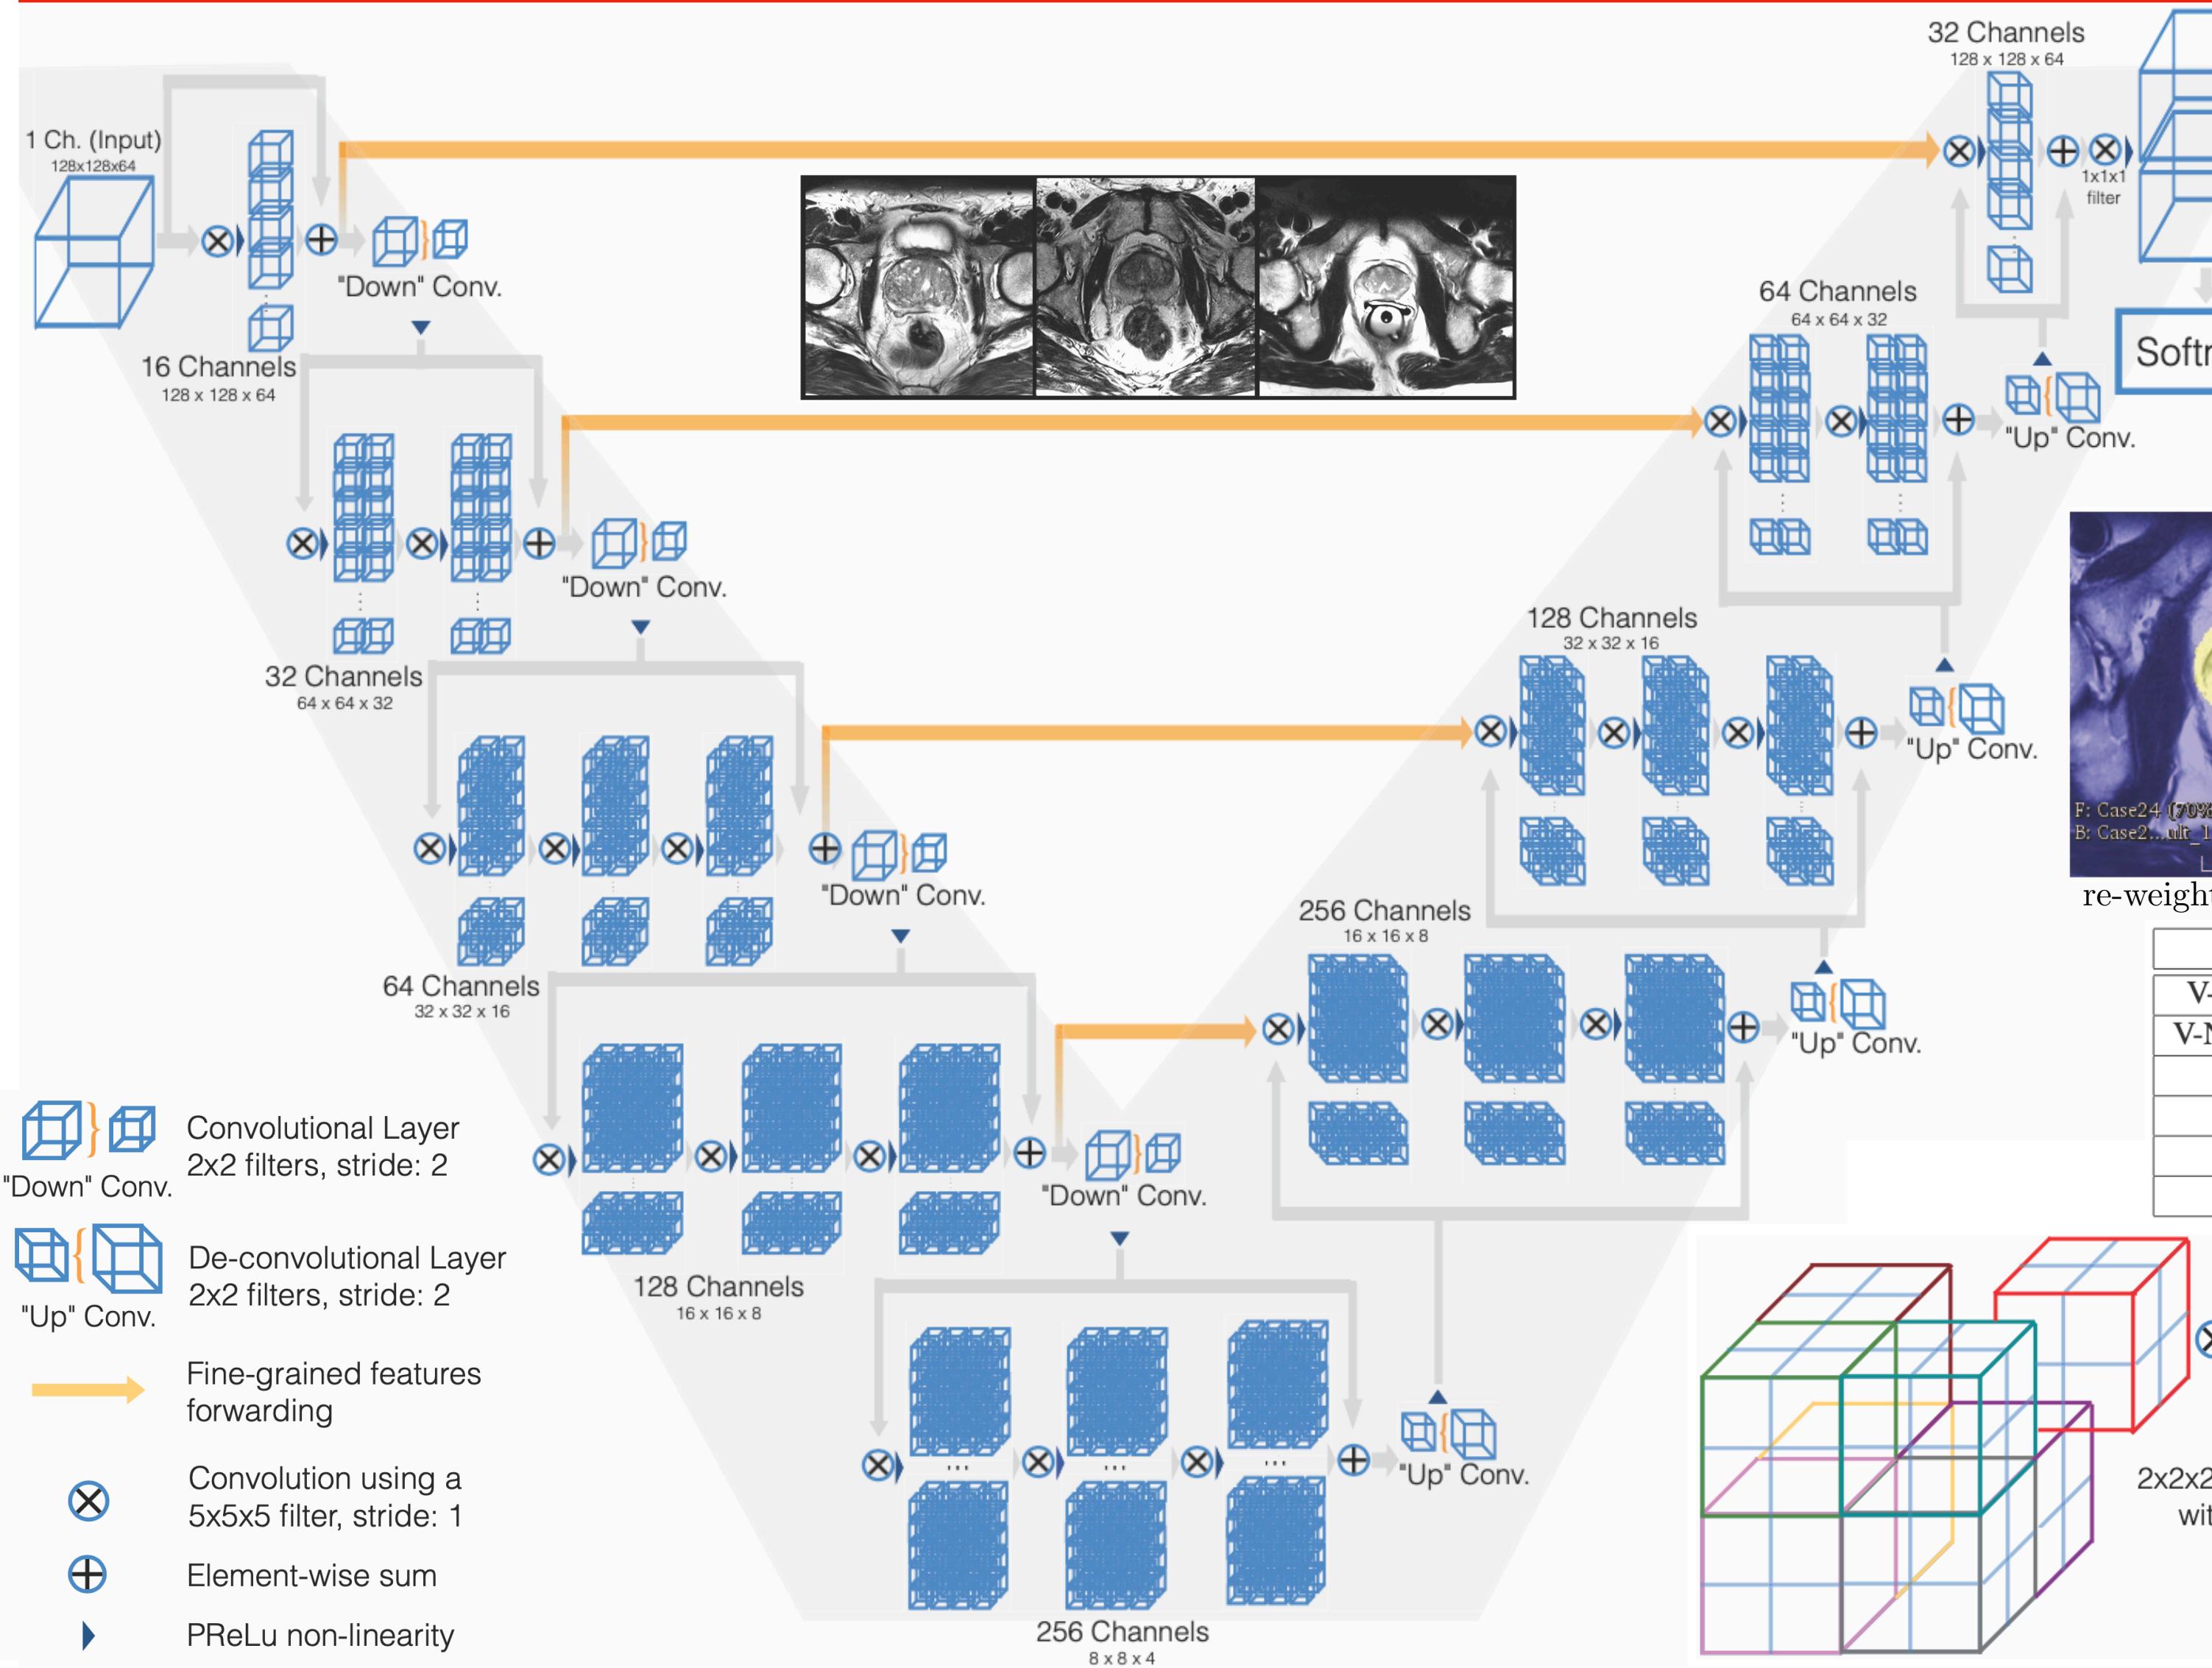
Department of Applied Mathematics

University of Colorado Boulder

maziar.raissi@colorado.edu



V-Net: Fully Convolutional Neural Networks for Volumetric Medical Image Segmentation



Dice Loss

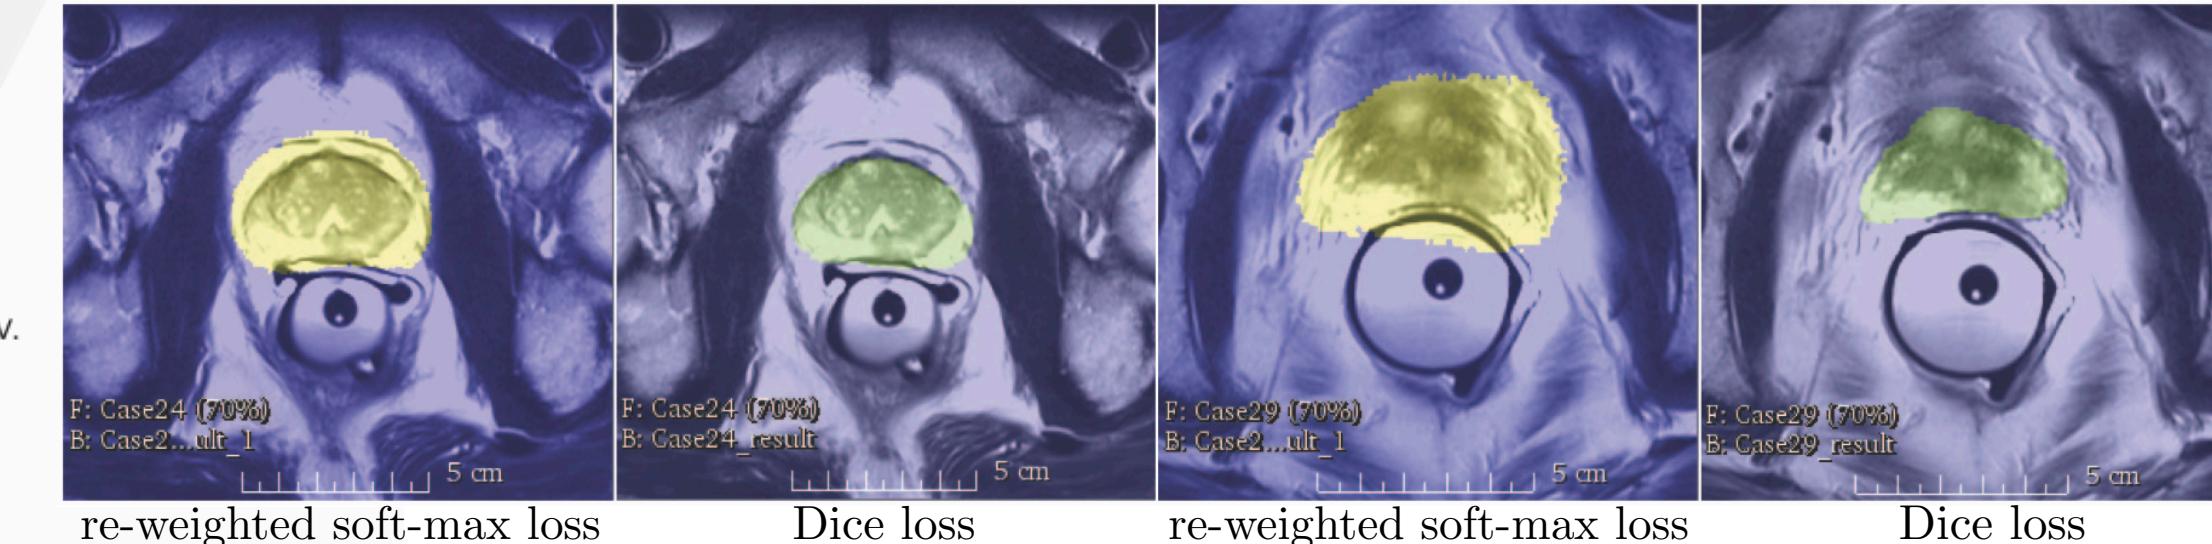
Class Imbalance! (Foreground v.s. Background)

$$D = \frac{2 \sum_i^N p_i g_i}{\sum_i^N p_i^2 + \sum_i^N g_i^2}$$

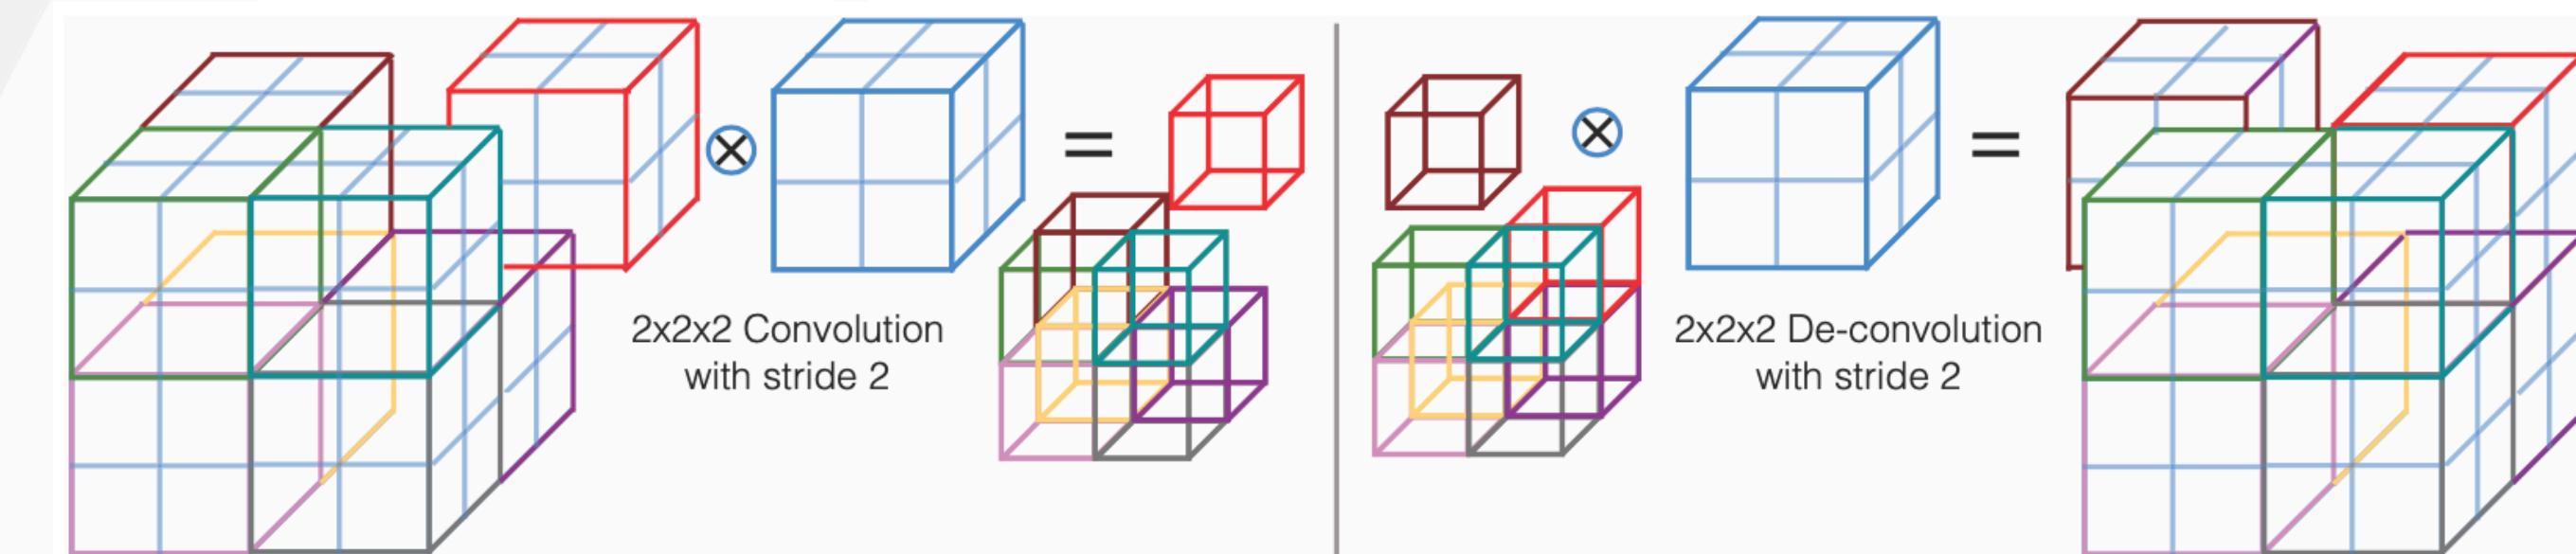
$N \rightarrow$ number of voxels

$p_i \in P \rightarrow$ predicted binary segmentation volume
 $g_i \in G \rightarrow$ ground truth binary volume

If $g_i = 0$, the background does not contribute to the loss!

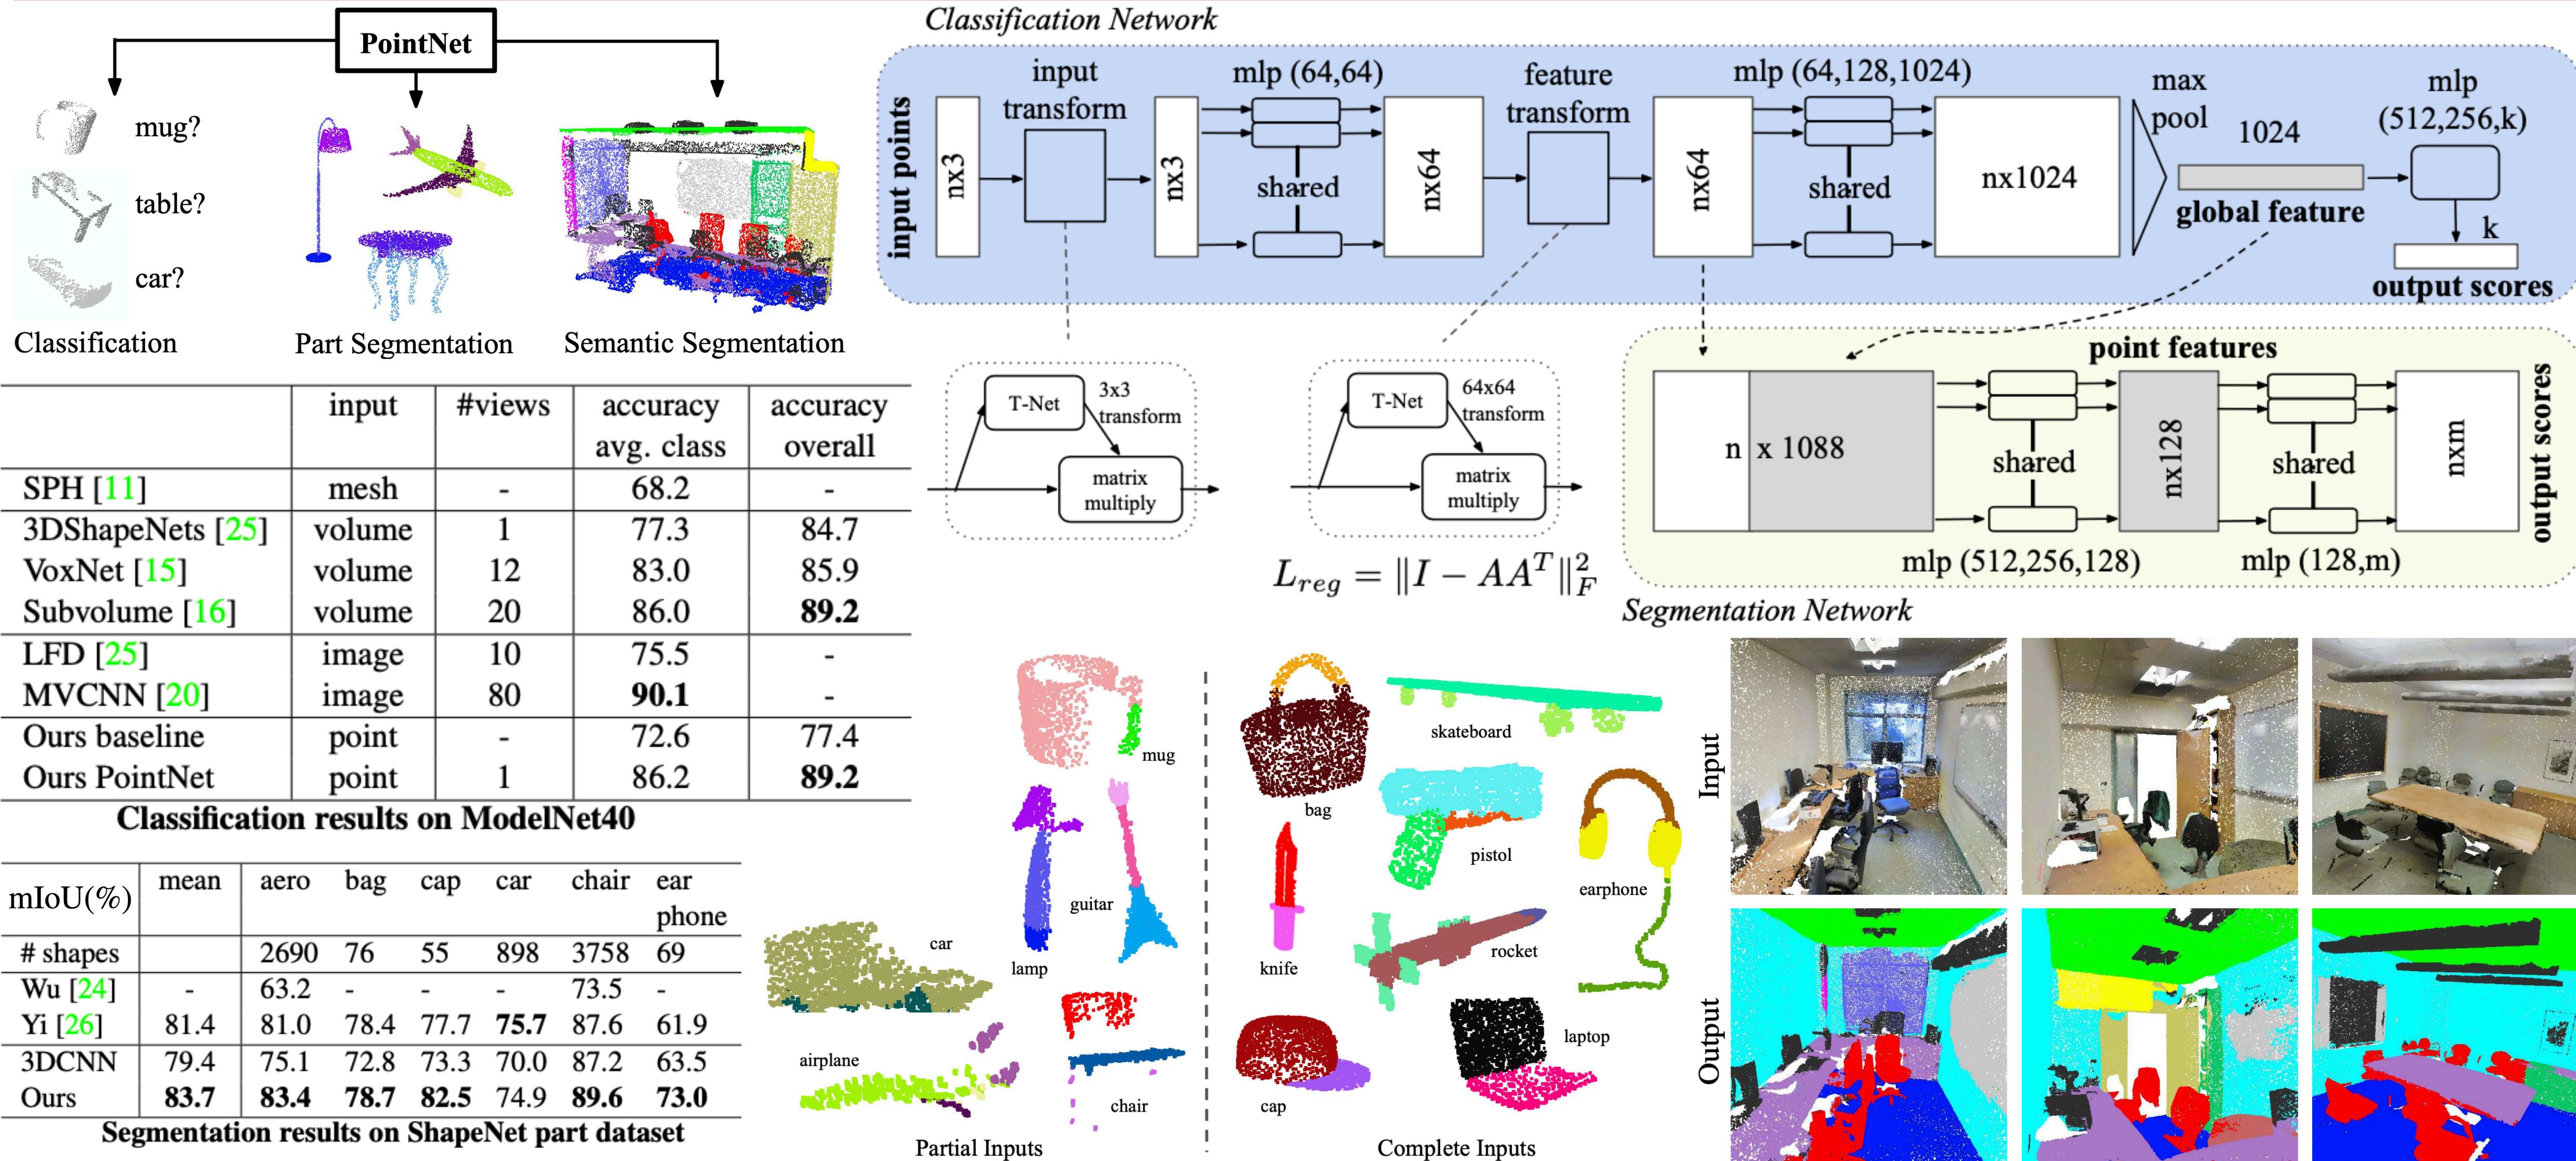


Algorithm	Avg. Dice	Avg. Hausdorff distance	Score on challenge task	Speed
V-Net + Dice-based loss	0.869 ± 0.033	$5.71 \pm 1.20 \text{ mm}$	82.39	1 sec.
V-Net + mult. logistic loss	0.739 ± 0.088	$10.55 \pm 5.38 \text{ mm}$	63.30	1 sec.
Imorphics [22]	0.879 ± 0.044	$5.935 \pm 2.14 \text{ mm}$	84.36	8 min.
ScrAutoProstate	0.874 ± 0.036	$5.58 \pm 1.49 \text{ mm}$	83.49	1 sec.
SBIA	0.835 ± 0.055	$7.73 \pm 2.68 \text{ mm}$	78.33	–
Grislies	0.834 ± 0.082	$7.90 \pm 3.82 \text{ mm}$	77.55	7 min.





PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation


[YouTube Video](#)


PointNet++: Deep Hierarchical Feature Learning on Point Sets in a Metric Space


[YouTube Video](#)

Visualization of a scan captured from a Structure Sensor



PointNet

$\{x_1, x_2, \dots, x_n\} \subset \mathbb{R}^d \rightarrow$ unordered point set
 $f(x_1, x_2, \dots, x_n) = \gamma(\max_{i=1, \dots, n}\{h(x_i)\})$
 $\gamma, h \rightarrow \text{MLP}$

Hierarchical Point Set Feature Learning

Local Context

- set abstraction levels
- sampling layer (centroids of local regions)
- grouping layer (finding “neighboring” points around centroids)
- PointNet layer (mini-PointNet)

Sampling Layer

Iterative Farthest Point Sampling (FPS)


$$\{x_{i_1}, x_{i_2}, \dots, x_{i_m}\} \subset \{x_1, x_2, \dots, x_n\}$$

$x_{i_j} \rightarrow$ the most distant point from the set $\{x_{i_1}, \dots, x_{i_{j-1}}\}$

Grouping Layer

$K \rightarrow$ number of points in the neighborhood of centroid points

Ball query works better than kNN!



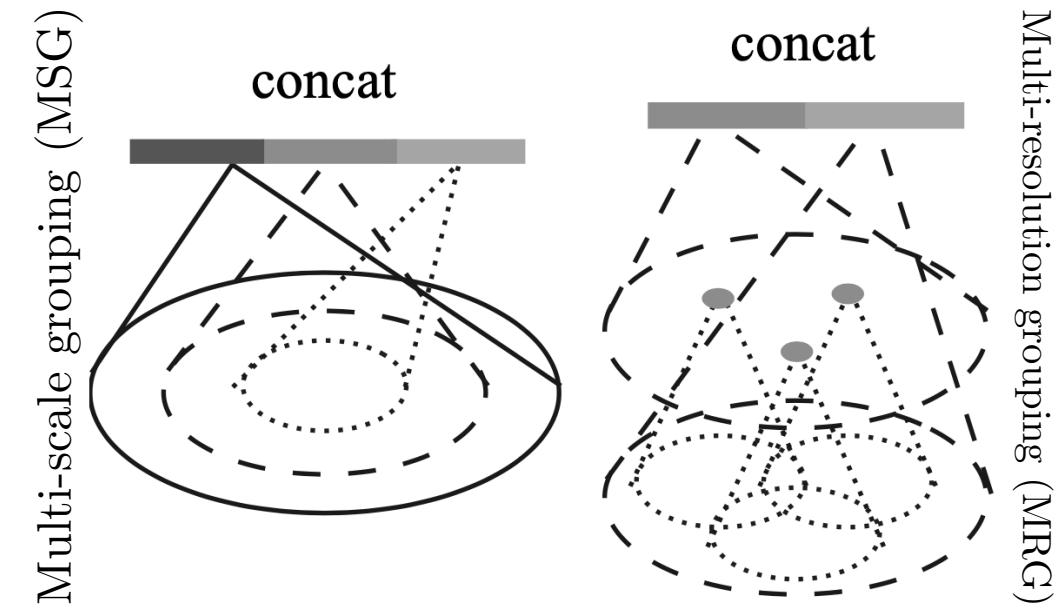
PointNet Layer

Relative coordinates

$$x_i^{(j)} \leftarrow x_i^{(j)} - \hat{x}^{(j)}, i = 1, \dots, K, j = 1, \dots, C$$

centroid

Non-Uniform Sampling Density



Random Input Dropout

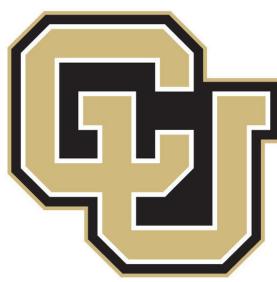
For each training point set, choose a dropout ratio $\theta \sim U[0, p]$ (e.g., $p = 0.95$). For each training point, randomly drop the point with probability θ .

Segmentation

$$f^{(j)}(x) = \frac{\sum_{i=1}^k w_i(x) f_i^{(j)}}{\sum_{i=1}^k w_i(x)}$$

$$\text{where } w_i(x) = \frac{1}{d(x, x_i)^p}, j = 1, \dots, C$$

Inverse Distance Weighted Interpolation averaged based on k nearest neighbors



Boulder

Dynamic Graph CNN for Learning on Point Clouds



[YouTube Video](#)

EdgeConv → edge convolution

- indoor navigation
- self-driving vehicles
- robotics
- shape synthesis & modeling

LiDAR scanners

- ModelNet40
- ShapeNetPart
- S3DIS

Edge Convolution

$$X = \{x_1, \dots, x_n\} \subset \mathbb{R}^F$$

$\underbrace{\quad}_{F\text{-dimensional point cloud with } n \text{ points}}$

$$F = 3 \rightarrow x_i = (x_i, y_i, z_i)$$

color, surface, normal, etc.

$\mathcal{G} = (\mathcal{V}, \mathcal{E}) \rightarrow$ directed graph representing local point cloud structure

$$\mathcal{V} = \{1, \dots, n\} \rightarrow \text{vertices}$$

$$\mathcal{E} \subset \mathcal{V} \times \mathcal{V} \rightarrow \text{edges}$$

k -nearest neighbor (k -NN) graph of X in \mathbb{R}^F

The graph includes self-loop

$$e_{ij} = h_\theta(x_i, x_j) \text{ where } h_\theta : \mathbb{R}^F \times \mathbb{R}^F \rightarrow \mathbb{R}^{F'}$$

$\underbrace{\quad}_{\text{edge feature}}$

$$x'_i = \square_{j:(i,j) \in \mathcal{E}} h_\theta(x_i, x_j) \quad \square = \sum \text{ or } \max$$

Choice of h & \square

$$x'_{im} = \sum_{j:(i,j) \in \mathcal{E}} \theta_m \cdot x_j \rightarrow \text{standard convolution}$$

$\underbrace{\quad}_{m\text{-th filter}}$

$$h_\theta(x_i, x_j) = h_\theta(x_i) \rightarrow \text{PointNet}$$

$$h_\theta(x_i, x_j) = h_\theta(x_j) \rightarrow \text{PointCNN}$$

$$x'_{im} = \sum_{j \in \mathcal{V}} h_\theta(x_j) g(u(x_i, x_j))$$

$\underbrace{\quad}_{\text{pairwise distance}}$

Gaussian kernel in Euclidean space

$$h_\theta(x_i, x_j) = h_\theta(x_i - x_j) \rightarrow \text{encode local information}$$

$$h_\theta(x_i, x_j) = \bar{h}_\theta(x_i, x_i - x_j) \rightarrow \text{encode global \& local information}$$

$$e'_{ijm} = \text{ReLU}(\theta_m \cdot (x_j - x_i) + \phi_m \cdot x_i)$$

$$x'_{im} = \max_{j:(i,j) \in \mathcal{E}} e'_{ijm}$$

Dynamic Graph CNN (DGCNN)

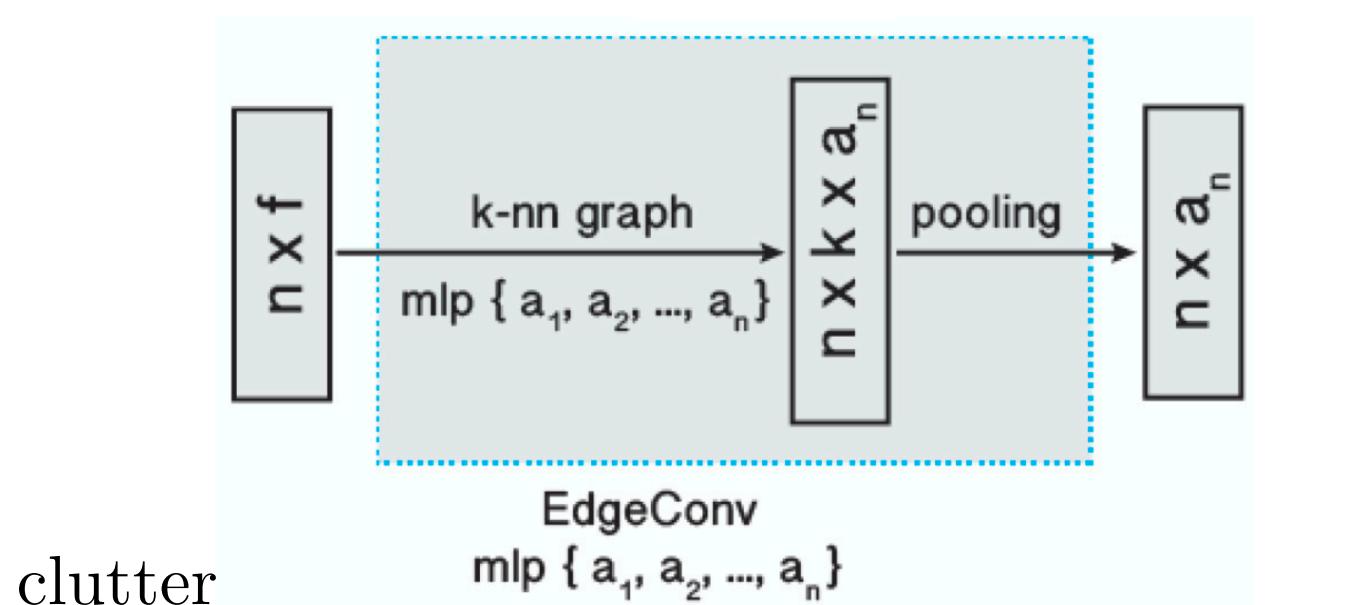

$\mathcal{G}^\ell = (\mathcal{V}^\ell, \mathcal{E}^\ell) \rightarrow$ dynamically constructed per each layer

Datasets

ModelNet40 contains 12,311 meshed CAD models from 40 categories.

ShapeNet part dataset contains 16,881 3D shapes from 16 object categories, annotated with 50 parts in total.

Stanford Large-Scale 3D Indoor Spaces Dataset (S3DIS) includes 3D scan point clouds for 6 indoor areas including 272 rooms in total. Each point belongs to one of 13 semantic categories—e.g., board, bookcase, chair, ceiling, and beam—plus clutter





Boulder



Questions?

[YouTube Playlist](#)
