



Boulder

# Generative Networks; Conditional GANs



[YouTube Playlist](#)

**Maziar Raissi**

**Assistant Professor**

Department of Applied Mathematics

University of Colorado Boulder

[maziar.raissi@colorado.edu](mailto:maziar.raissi@colorado.edu)



Boulder



[YouTube Video](#)

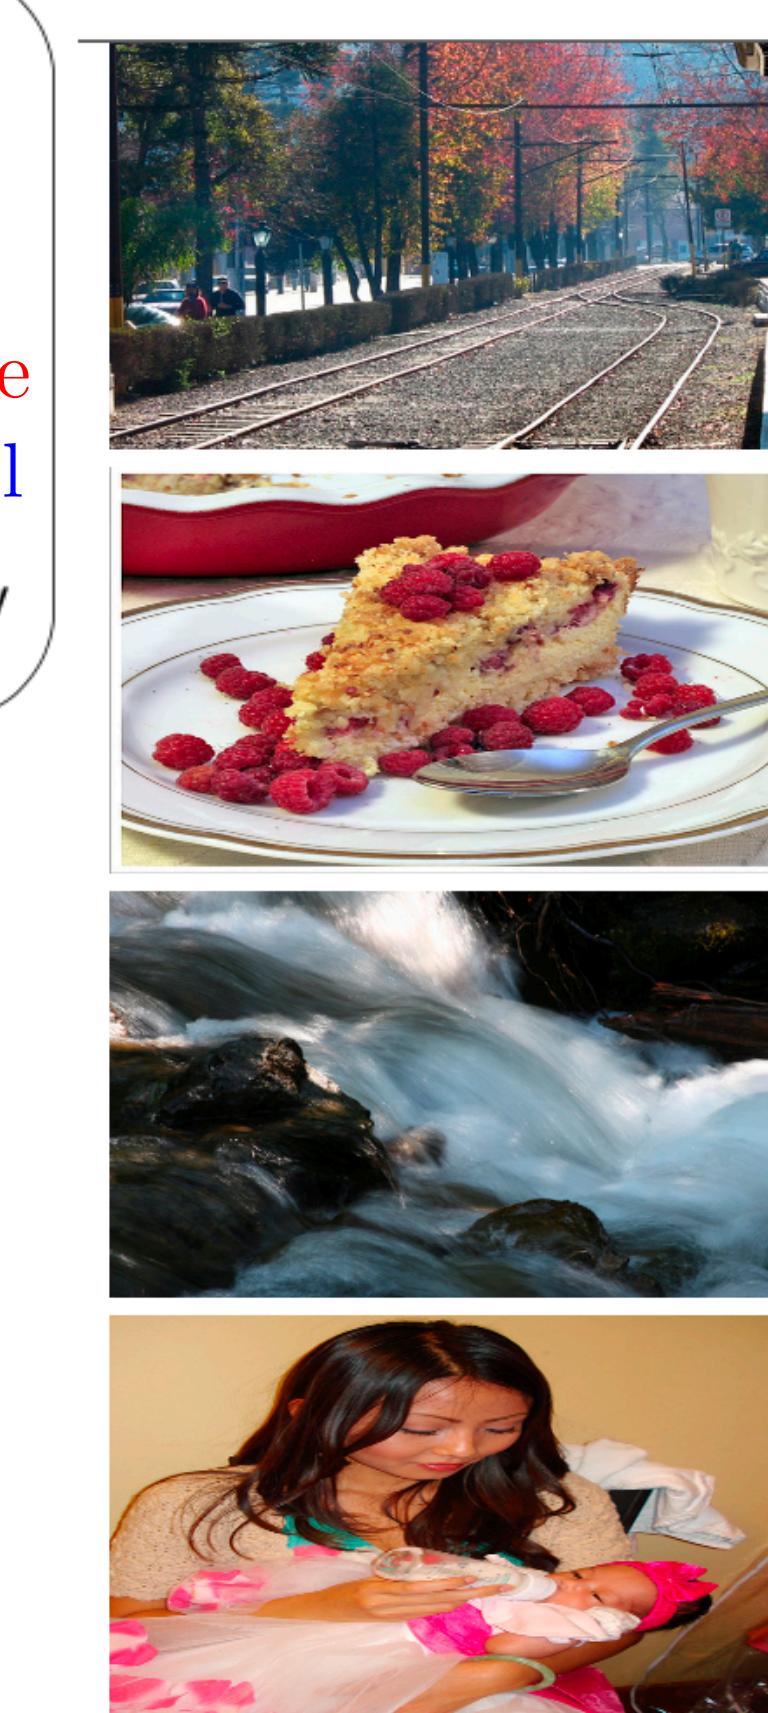
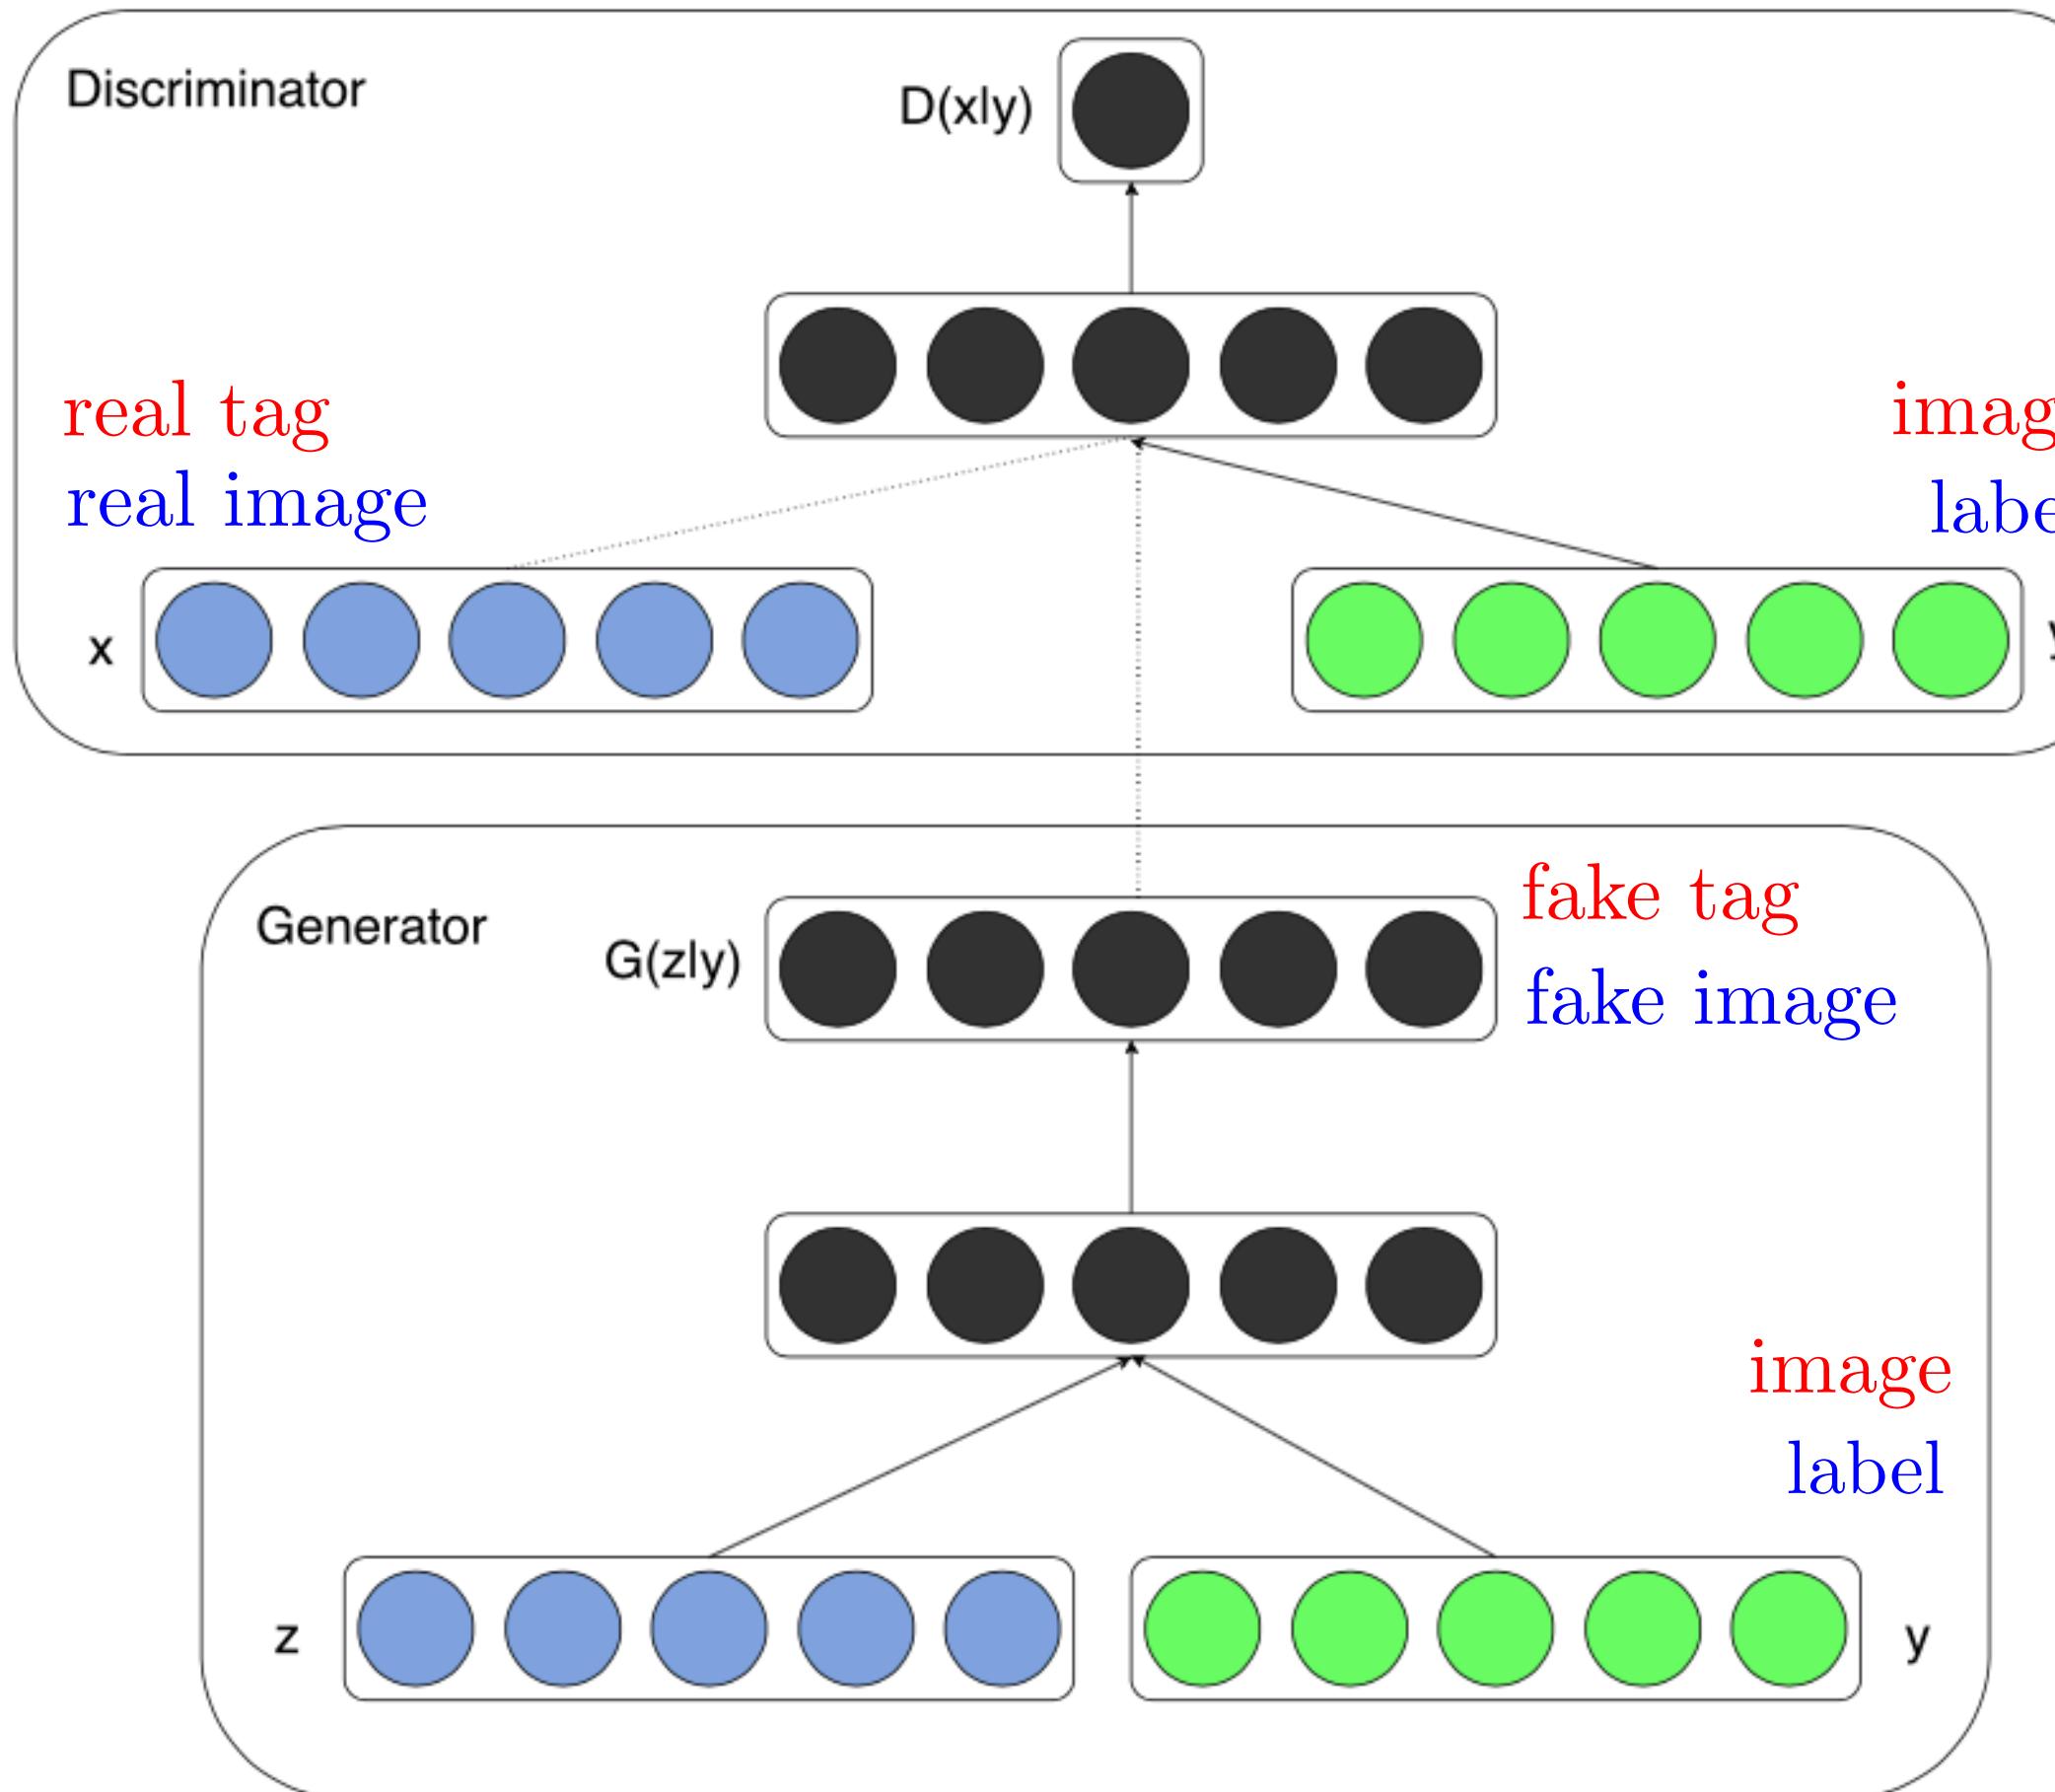
# Conditional Generative Adversarial Nets

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]$$

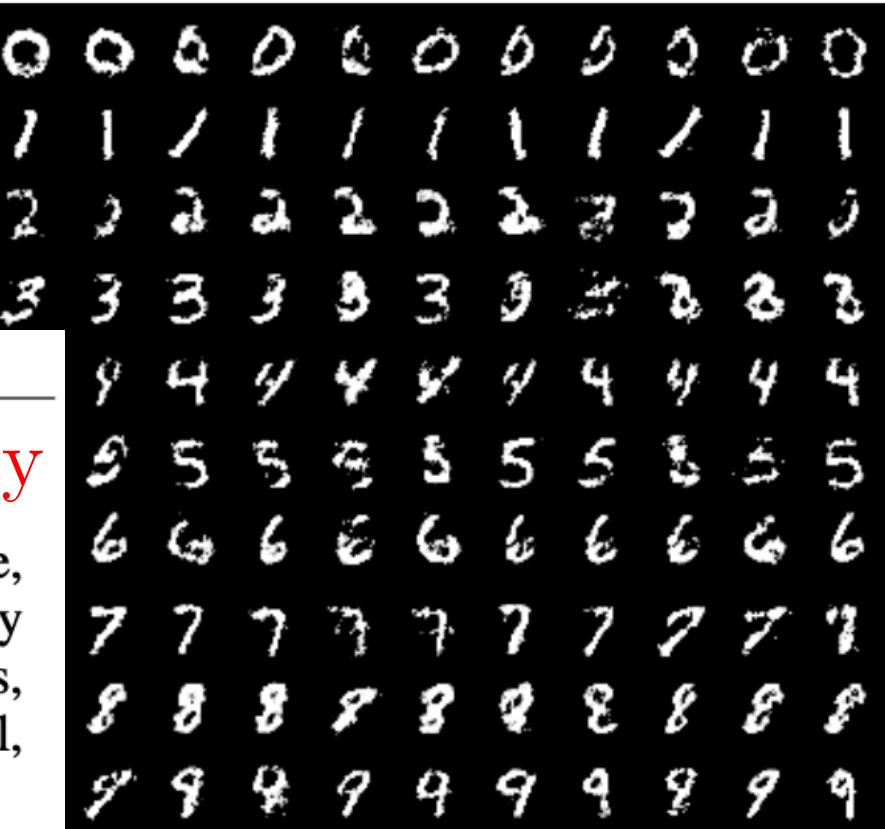
$D(x|y)$

$G(z|y)$

One-To-Many Mapping



User tags + annotations	Generated tags
montanha, trem, inverno, frio, people, male, plant life, tree, structures, transport, car	cosine similarity 9 4 4 4 4 4 4 4 4 4 9 5 5 5 5 5 5 5 5 5 6 6 6 6 6 6 6 6 6 6 7 7 7 7 7 7 7 7 7 7 8 8 8 8 8 8 8 8 8 8 9 9 9 9 9 9 9 9 9 9
food, raspberry, delicious, homemade	chicken, fattening, cooked, peanut, cream, cookie, house made, bread, biscuit, bakes
water, river	creek, lake, along, near, river, rocky, treeline, valley, woods, waters
people, portrait, female, baby, indoor	love, people, posing, girl, young, strangers, pretty, women, happy, life



tag: word vector  
label: one-hot vector



Boulder



[YouTube Video](#)

# Context Encoders: Feature Learning by Inpainting



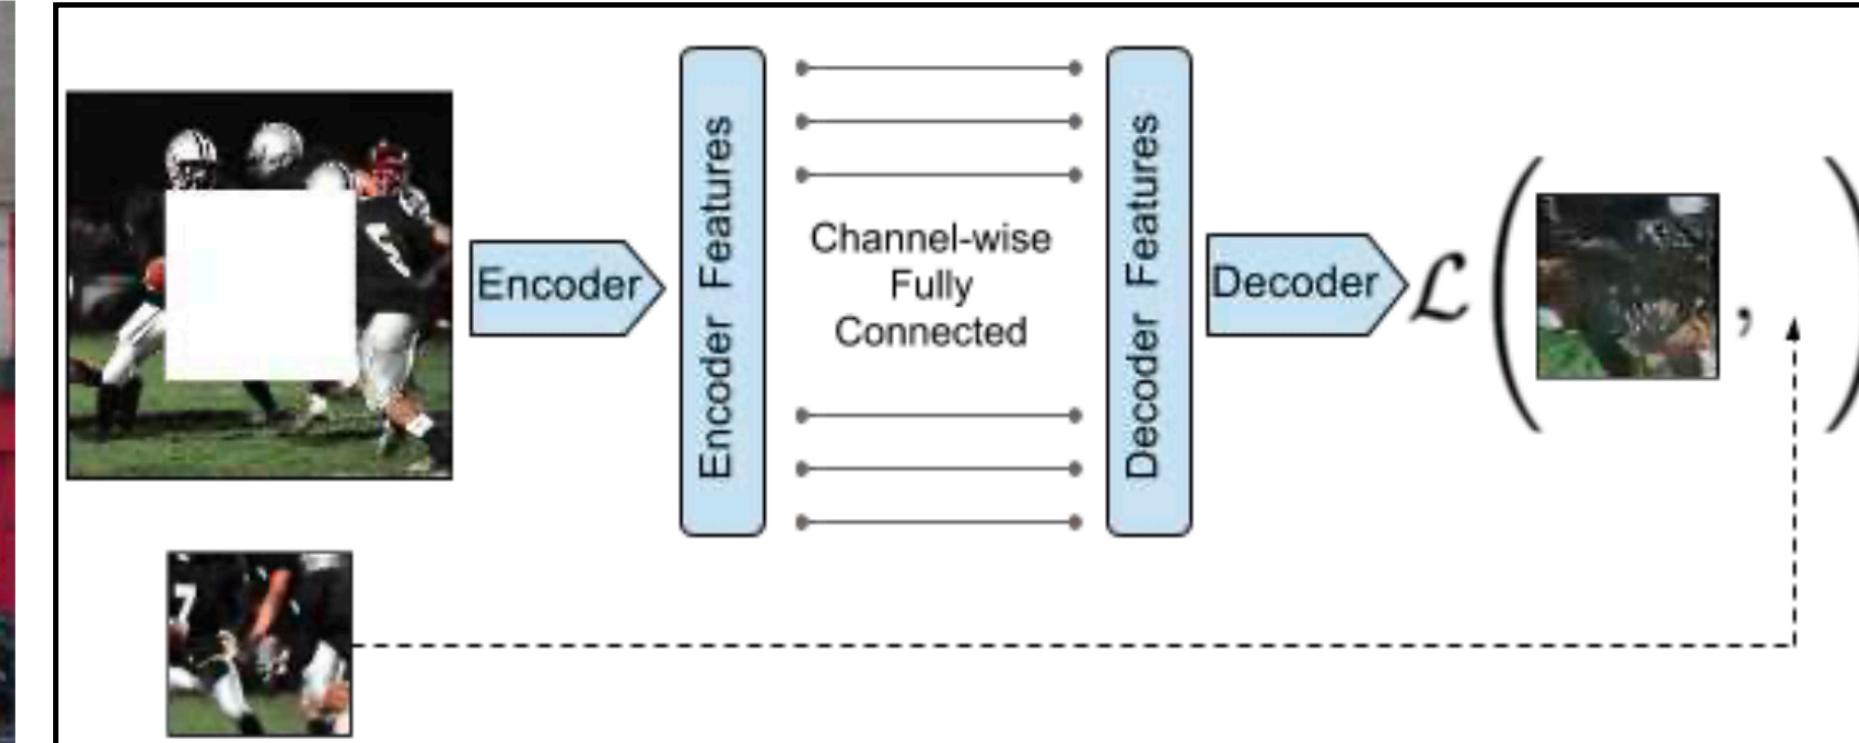
(a) Input context



(c) Context Encoder  
(L2 loss)

**Loss function**  
 $x \rightarrow$  ground truth image  
 $F \rightarrow$  context encoder

$F(x) \rightarrow$  output  
 $\hat{M} \rightarrow$  binary mask (dropped image region)  
 1 whenever a pixel is dropped and 0 for input pixels



**Encoder**  
 $\mathbb{R}^{227 \times 227 \times 3} \ni I \triangleright \underbrace{\text{AlexNet(pool5)}}_{\substack{5 \text{ conv layers} \\ 6 \times 6 \times 256 = 9,216 \text{ from scratch}} \in \mathbb{R}^{6 \times 6 \times 256}}$

**Channel-wise fully-connected layer**

input:  $m$  feature maps of size  $n \times n$

output:  $m$  feature maps of size  $n \times n$

$$Y_j = \underbrace{W_j}_{n^2 \times 1} \underbrace{X_j}_{n^2 \times n^2} \in \mathbb{R}^{n^2 \times 1}$$

$mn^4 \rightarrow$  number of parameters (rather than  $m^2n^4$ )

**Decoder**

5 up-conv layers

$$\mathcal{L}_{rec}(x) = \|\hat{M} \odot (x - F((1 - \hat{M}) \odot x))\|_2$$

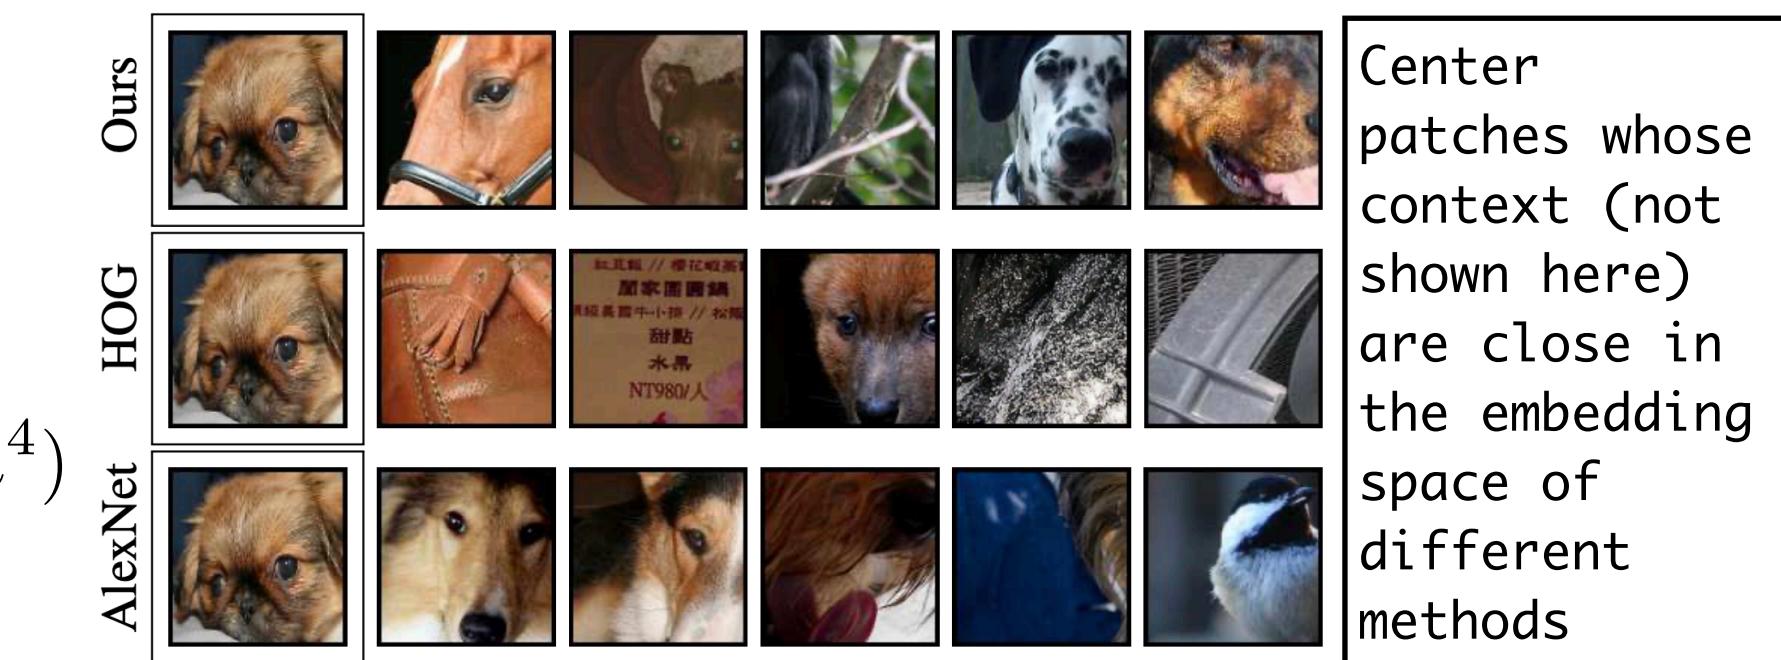
$$\mathcal{L}_{adv} = \max_D \mathbb{E}_{x \in \mathcal{X}} [\log(D(x))]$$

$$+ \log(1 - D(F((1 - \hat{M}) \odot x)))]$$

$$\mathcal{L} = \lambda_{rec}\mathcal{L}_{rec} + \lambda_{adv}\mathcal{L}_{adv}$$



Method	Mean L1 Loss	Mean L2 Loss	PSNR (higher better)
NN-inpainting (HOG features)	19.92%	6.92%	12.79 dB
NN-inpainting (our features)	15.10%	4.30%	14.70 dB
Our Reconstruction (joint)	<b>10.33%</b>	<b>2.35%</b>	<b>17.59 dB</b>



Pretraining Method	Supervision	Pretraining time	Classification	Detection	Segmentation
ImageNet [26]	1000 class labels	3 days	<b>78.2%</b>	<b>56.8%</b>	<b>48.0%</b>
Random Gaussian	initialization	< 1 minute	53.3%	43.4%	19.8%
Autoencoder	-	14 hours	53.8%	41.9%	25.2%
Agrawal <i>et al.</i> [1]	egomotion	10 hours	52.9%	41.8%	-
Doersh <i>et al.</i> [7]	context	4 weeks	55.3%	<b>46.6%</b>	-
Wang <i>et al.</i> [39]	motion	1 week	<b>58.4%</b>	44.0%	-
Ours	context	14 hours	56.5%	44.5%	<b>29.7%</b>

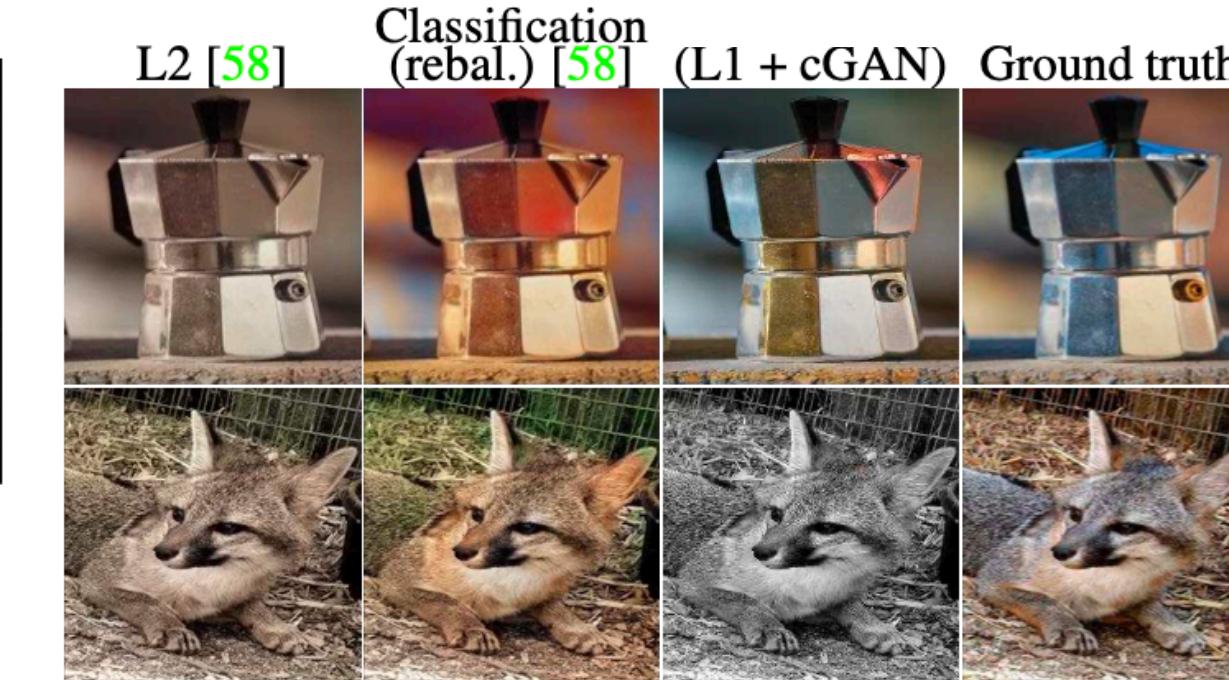
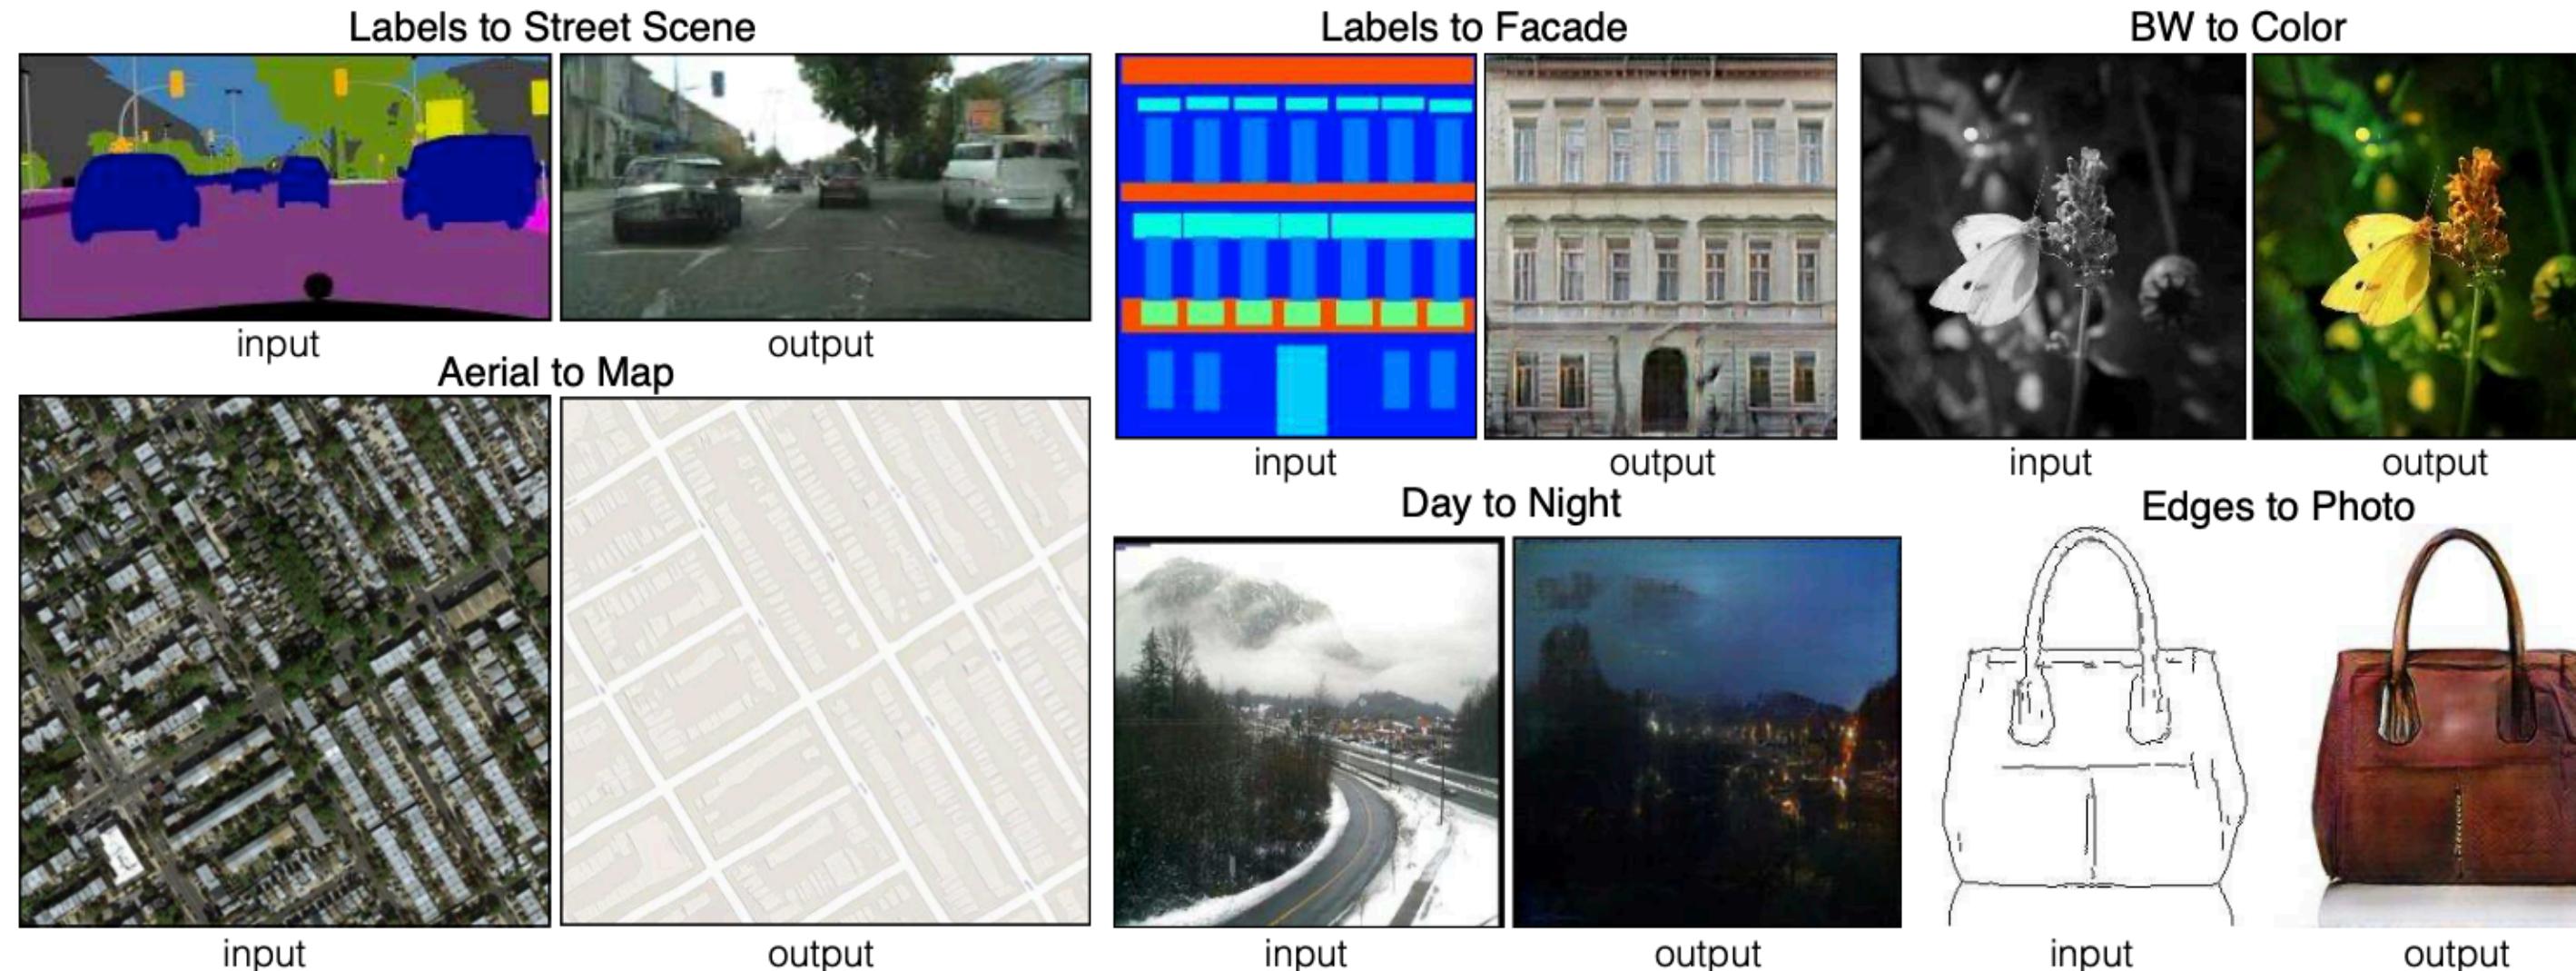


Boulder

# Image-to-Image Translation with Conditional Adversarial Networks



[YouTube Video](#)



Background removal



by Kaihu Chen

“Do as I do”



by Brannon Dorsey

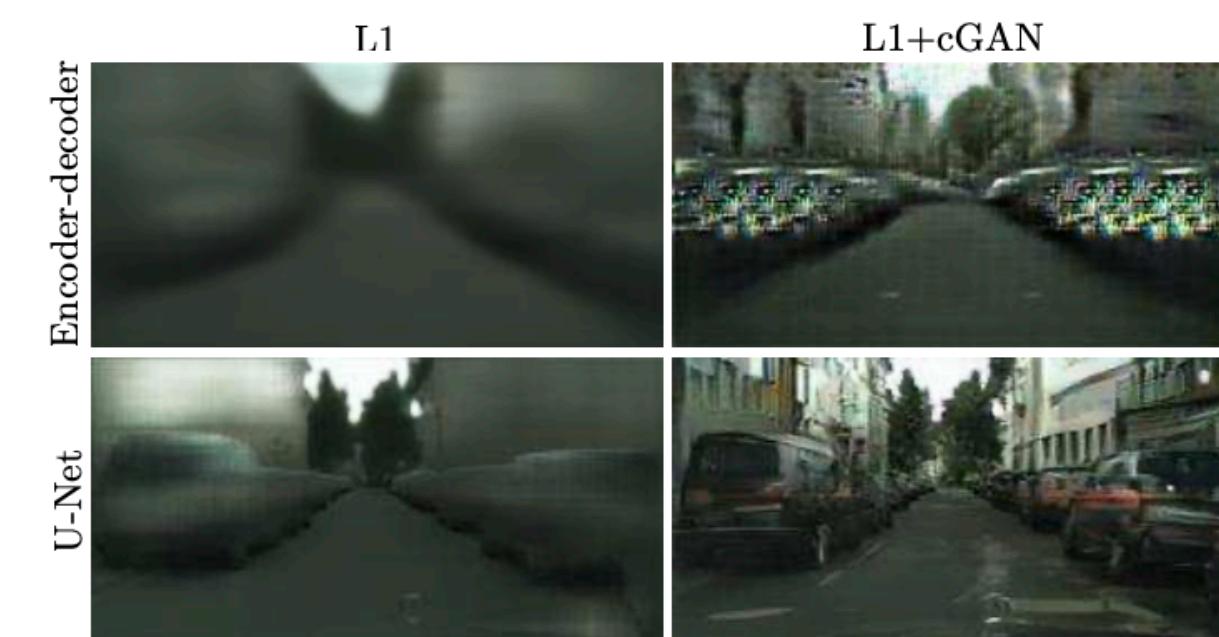
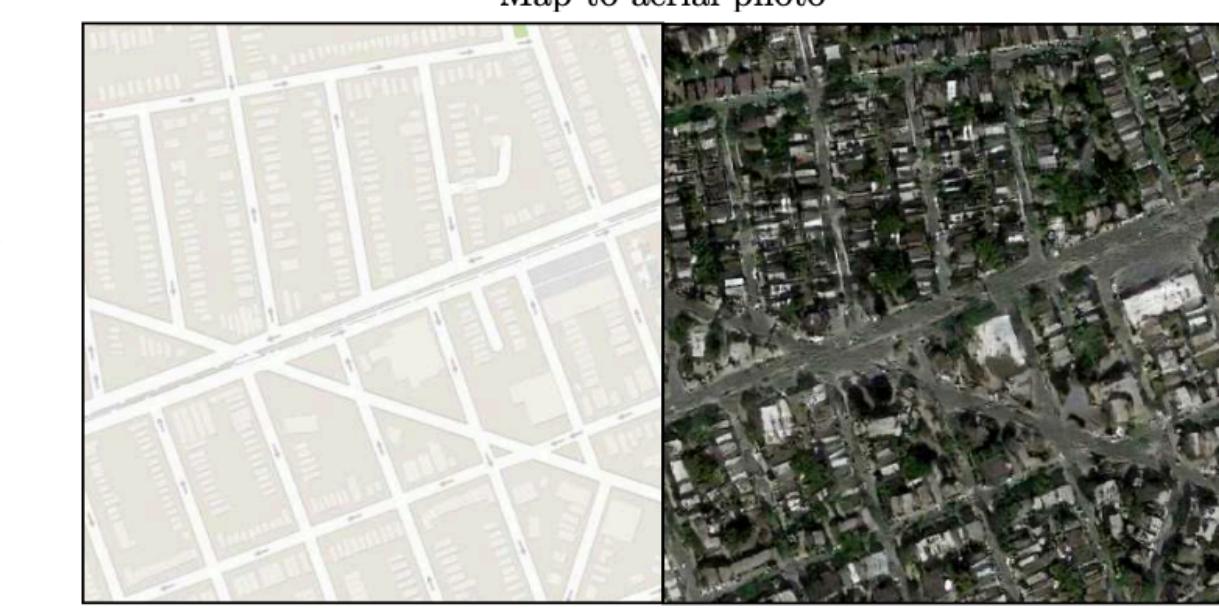
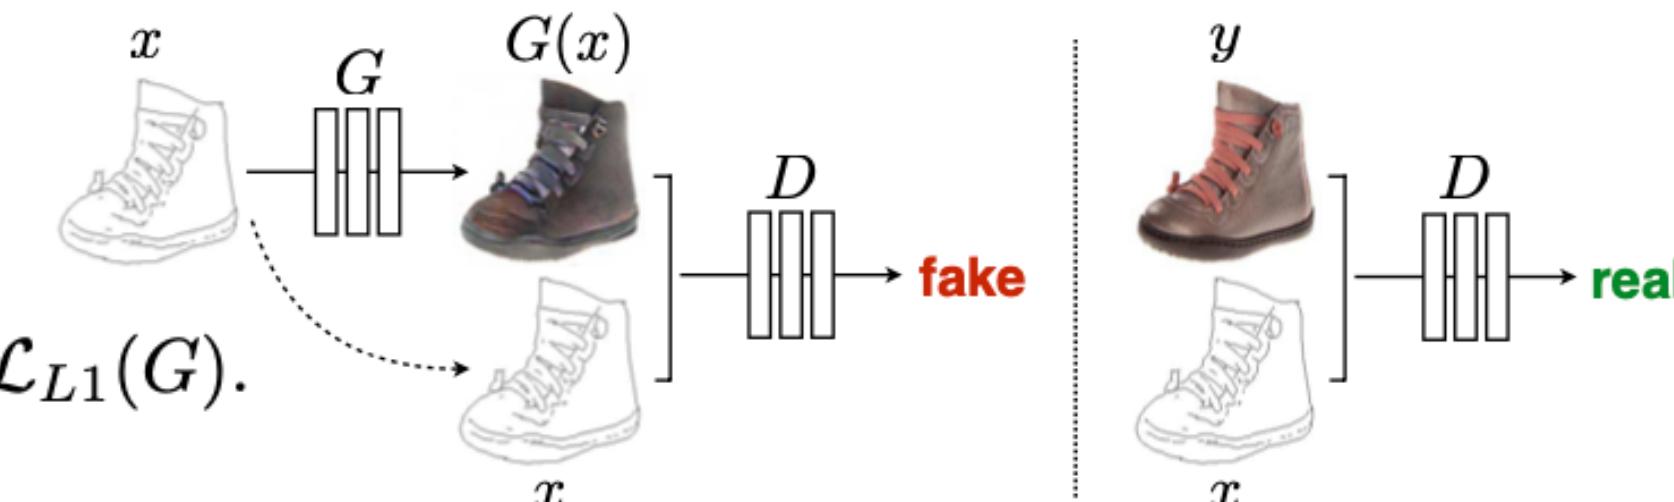
$$\mathcal{L}_{cGAN}(G, D) = \mathbb{E}_{x,y}[\log D(x, y)] + \\ \mathbb{E}_{x,z}[\log(1 - D(x, G(x, z)))] \rightarrow \text{objective of a conditional GAN}$$

$$G^* = \arg \min_G \max_D \mathcal{L}_{cGAN}(G, D)$$

$$\mathcal{L}_{L1}(G) = \mathbb{E}_{x,y,z}[\|y - G(x, z)\|_1].$$

$$G^* = \arg \min_G \max_D \mathcal{L}_{cGAN}(G, D) + \lambda \mathcal{L}_{L1}(G).$$

$$\mathcal{L}_{GAN}(G, D) = \mathbb{E}_y[\log D(y)] + \\ \mathbb{E}_{x,z}[\log(1 - D(G(x, z)))] \rightarrow \text{an unconditional variant}$$



Isola, Phillip, et al. "Image-to-image translation with conditional adversarial networks." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2017.

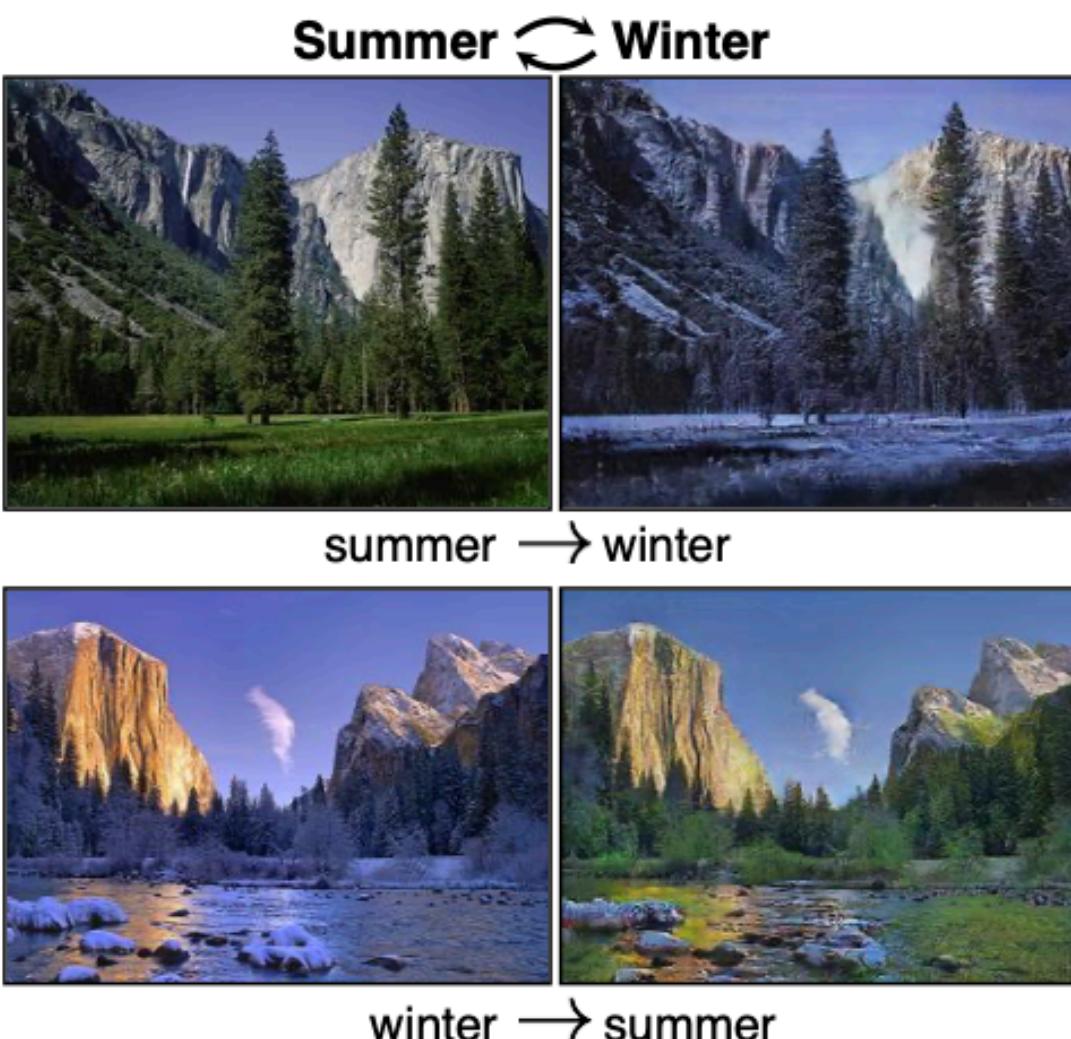
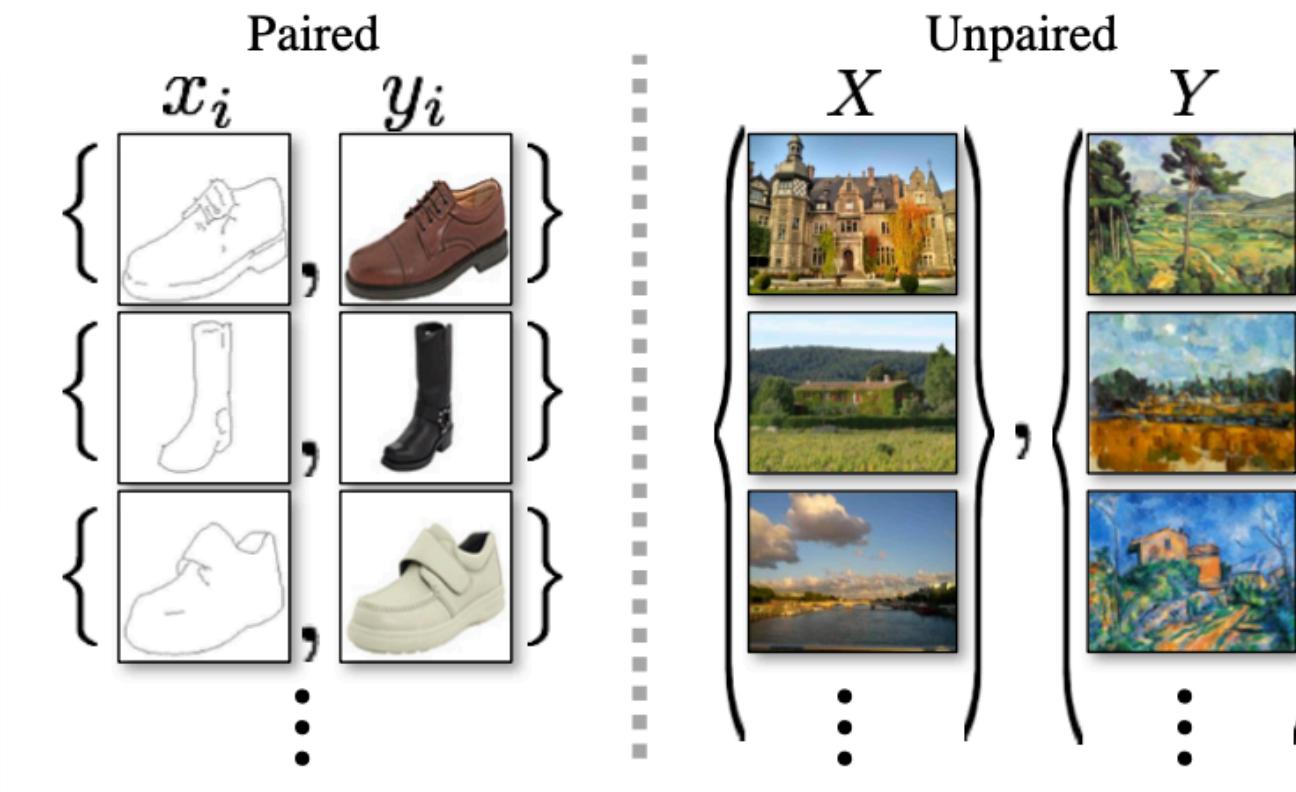
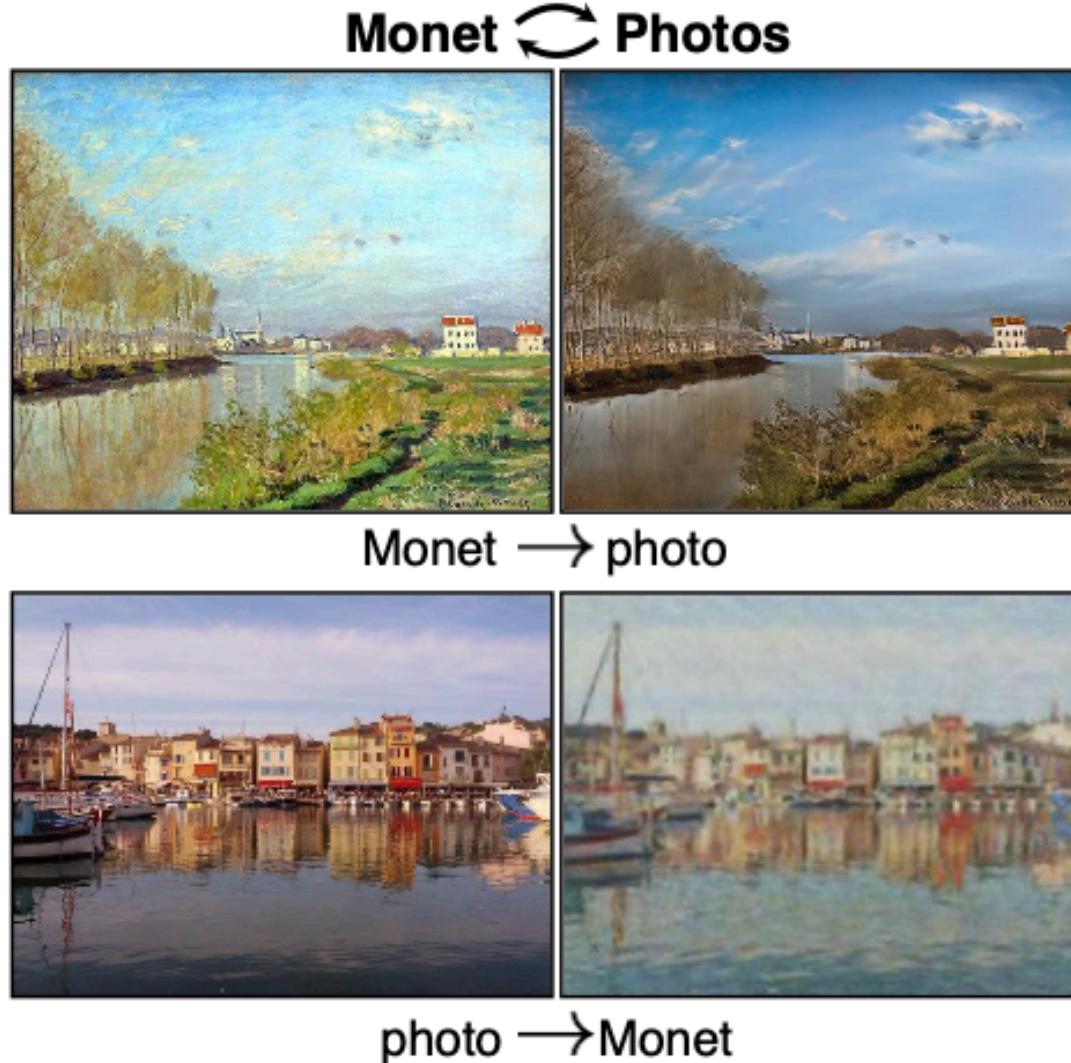


Boulder

# Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks



[YouTube Video](#)



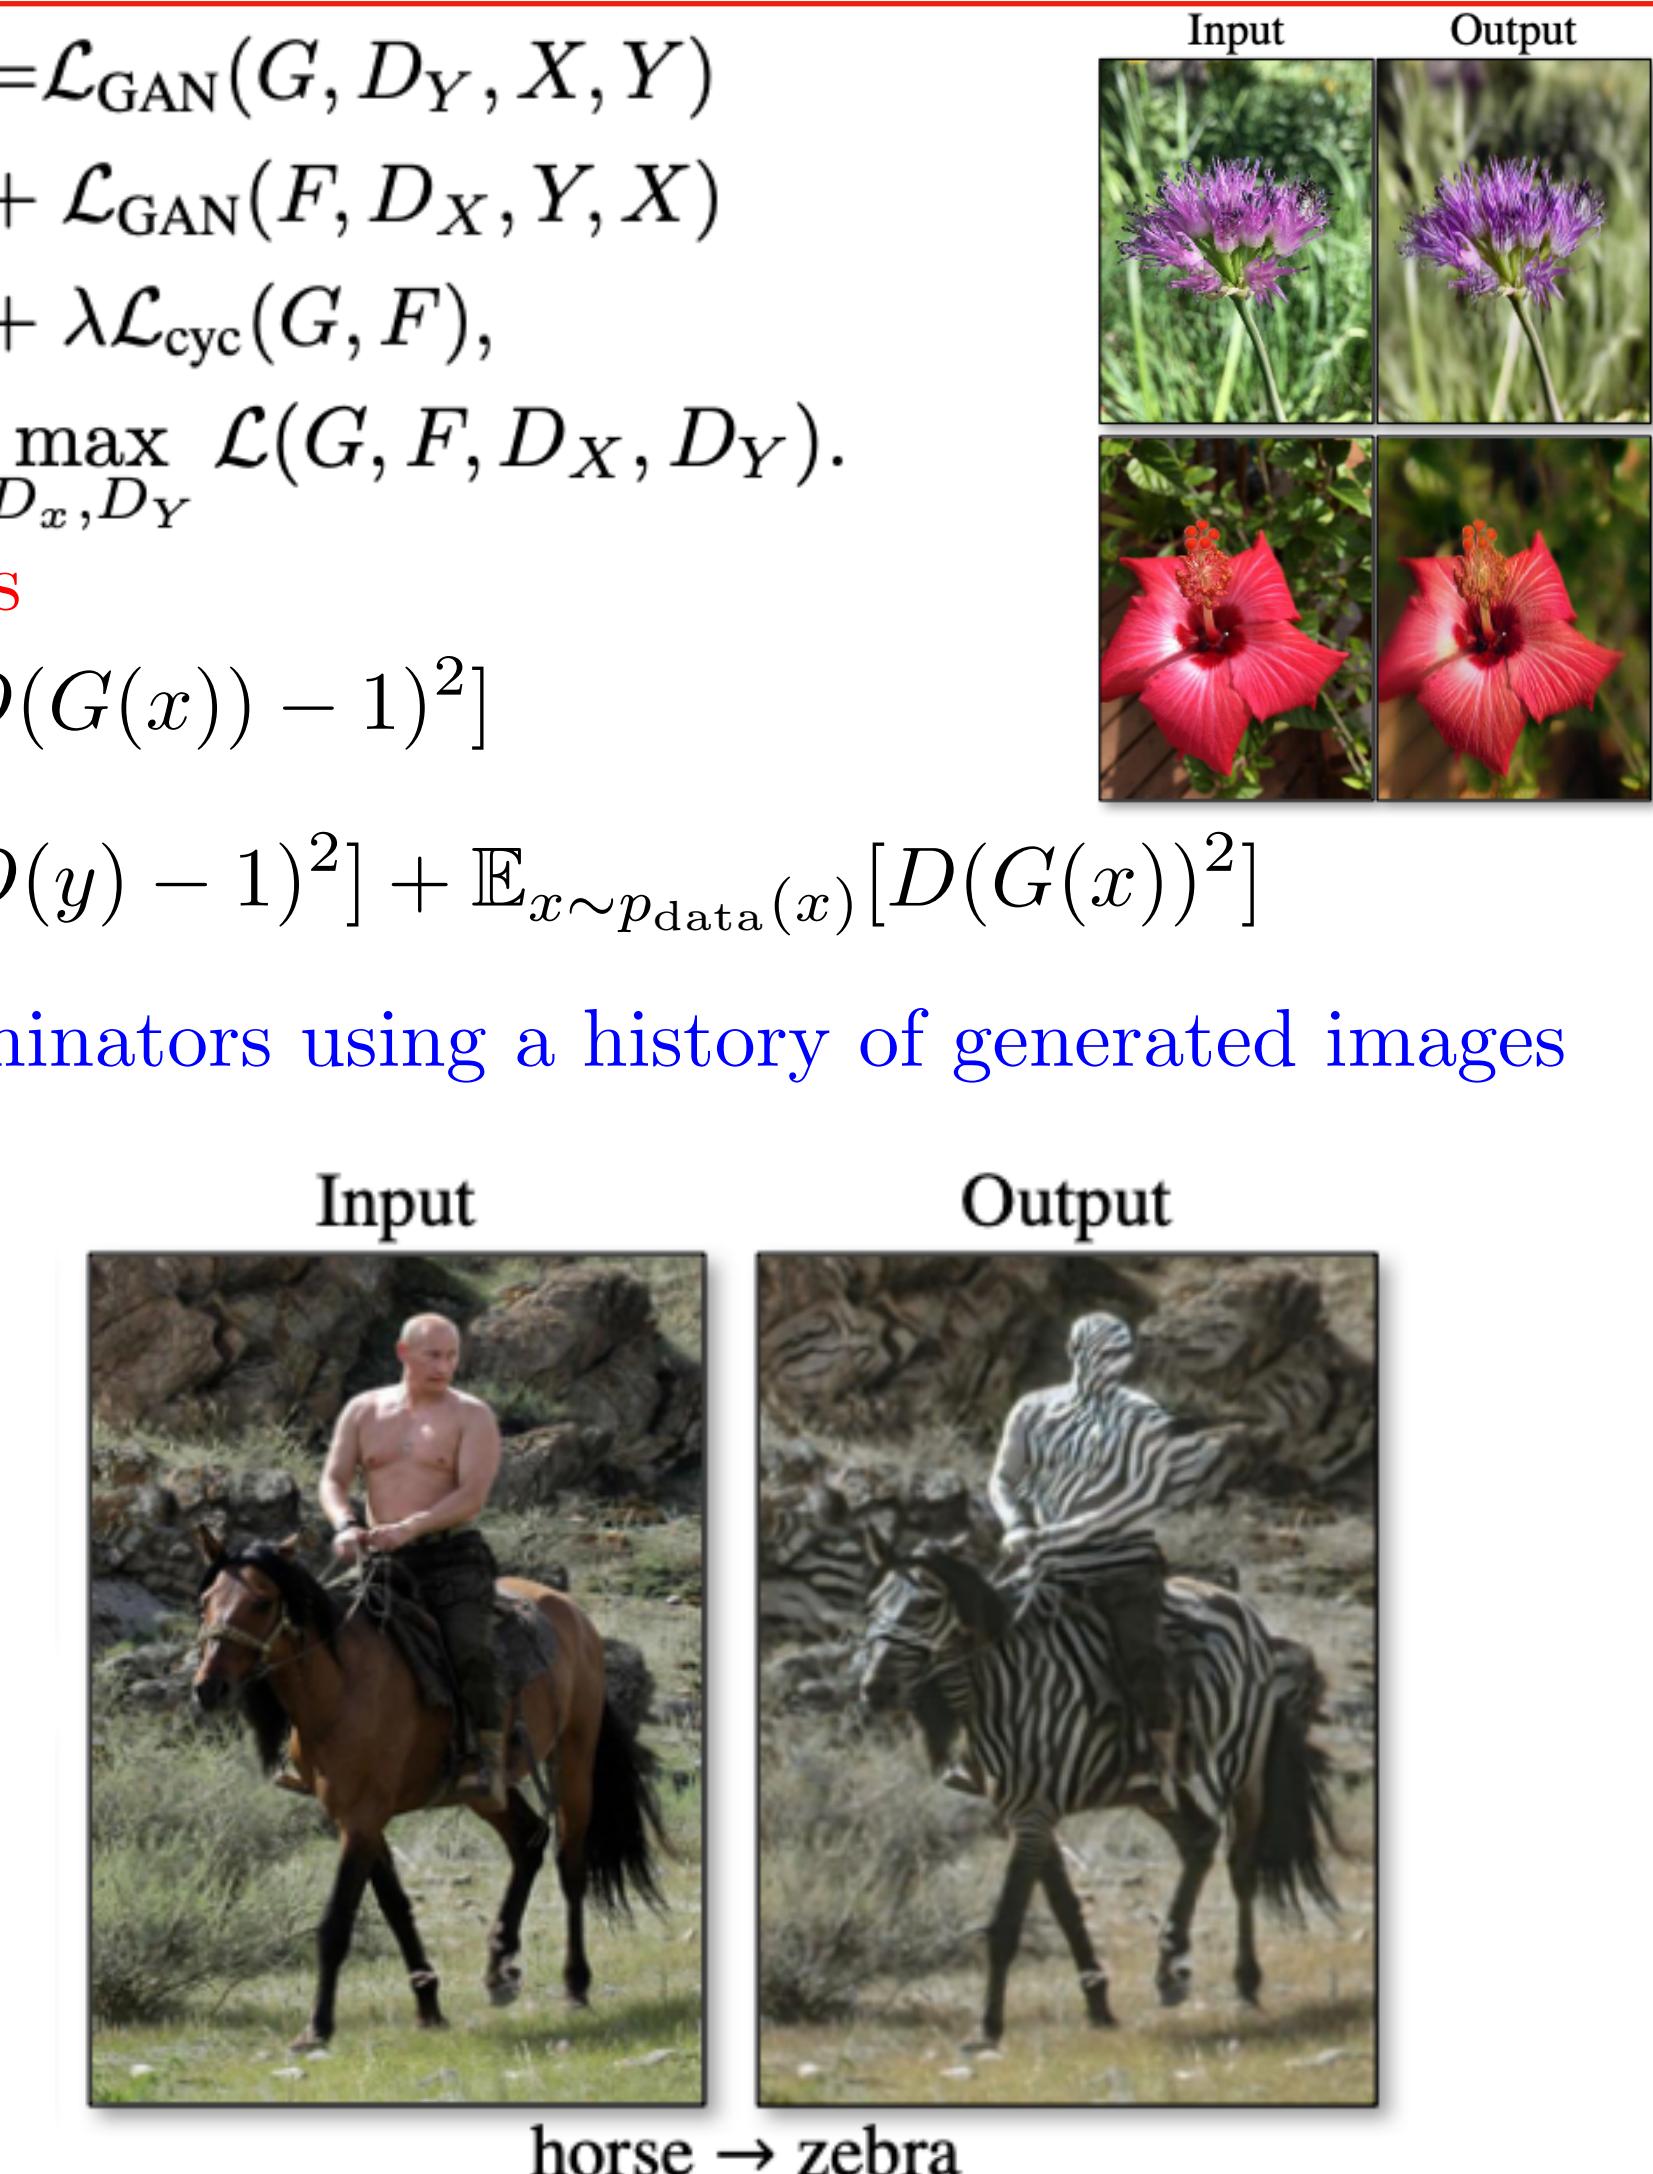
$$\begin{aligned} \mathcal{L}(G, F, D_X, D_Y) = & \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) \\ & + \mathcal{L}_{\text{GAN}}(F, D_X, Y, X) \\ & + \lambda \mathcal{L}_{\text{cyc}}(G, F), \\ G^*, F^* = & \arg \min_{G, F} \max_{D_x, D_Y} \mathcal{L}(G, F, D_X, D_Y). \end{aligned}$$

Least-squares Loss

$$\min_G \mathbb{E}_{x \sim p_{\text{data}}(x)} [(D(G(x)) - 1)^2]$$

$$\min_D \mathbb{E}_{y \sim p_{\text{data}}(y)} [(D(y) - 1)^2] + \mathbb{E}_{x \sim p_{\text{data}}(x)} [D(G(x))^2]$$

update the discriminators using a history of generated images





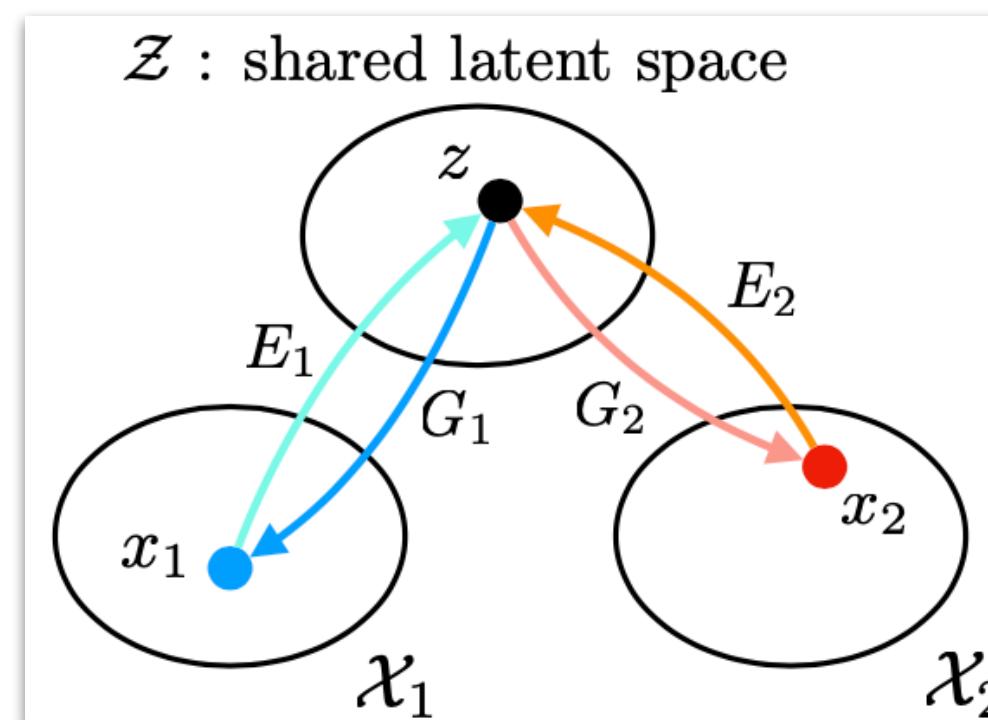
Boulder

# Unsupervised Image-to-Image Translation Networks

## UNIT: UNsupervised Image-to-image Translation Networks

Goal: To learn a joint distribution of images in different domains by using images from the marginal distributions in individual domains.  
Since an infinite set of possible joint distributions can yield the given marginal distributions, we could infer nothing about the joint distribution from the marginal samples without additional assumptions.

### Shared Latent Space Assumption



$$z = E_1(x_1) = E_2(x_2)$$

$$x_1 = G_1(z) \quad x_2 = G_2(z)$$

Shared latent space constraint implies cycle-consistency!

$$F^{1 \rightarrow 2} := G_2 \circ E_1 \quad F^{2 \rightarrow 1} := G_1 \circ E_2$$

$$x_1 = F^{2 \rightarrow 1}(F^{1 \rightarrow 2}(x_1))$$

$$x_2 = F^{1 \rightarrow 2}(F^{2 \rightarrow 1}(x_2))$$

$\{E_1, E_2\} \rightarrow$  encoders,  $\{G_1, G_2\} \rightarrow$  generators,  $\{D_1, D_2\} \rightarrow$  discriminators

### Variational Autoencoders (VAEs)

$\{E_1, G_1\} \rightarrow$  VAE for the  $\mathcal{X}_1$  domain ( $\text{VAE}_1$ )

$$q_1(z_1|x_1) := \mathcal{N}(z_1|E_1(x_1), I)$$

$x_1^{1 \rightarrow 1} = G_1(\underbrace{z_1 \sim q_1(z_1|x_1)}_{\text{re-parametrization trick}}) \rightarrow$  reconstructed image

$$\text{re-parametrization trick } z_1 = E_1(x_1) + \eta, \eta \sim \mathcal{N}(0, I)$$

$$\mathcal{L}_{\text{VAE}_1}(E_1, G_1) = \lambda_1 \text{KL}(q_1(z_1|x_1)||p_\eta(z)) - \lambda_2 \mathbb{E}_{z_1 \sim q_1(z_1|x_1)} [\log p_{G_1}(x_1|z_1)]$$

$$p_\eta(z) = \mathcal{N}(z|0, I)$$

$p_{G_1} \rightarrow$  Laplace distribution (equivalent to minimizing the absolute distance)

### Weight Sharing

$$E_1 = E_H \circ E_{L,1}, E_2 = E_H \circ E_{L,2}$$

$H, L \rightarrow$  High, Low level features

$$G_1 = G_{L,1} \circ G_H, G_2 = G_{L,2} \circ G_H$$

Note that the weight-sharing constraint alone does not guarantee that corresponding images in two domains will have the same latent code.

### Generative Adversarial Networks (GANs)

$$\{G_1, D_1\} \rightarrow \text{GAN}_1$$

$G_1$  can generate two types of images:

- reconstruction stream (supervised training):  $x_1^{1 \rightarrow 1} = G_1(z_1 \sim q_1(z_1|x_1))$
- translation stream (adversarial training):  $x_2^{2 \rightarrow 1} = G_1(z_2 \sim q_2(z_2|x_2))$

$$\mathcal{L}_{\text{GAN}_1}(E_1, G_1, D_1) = \lambda_0 \mathbb{E}_{x_1 \sim P_{\mathcal{X}_1}} [\log D_1(x_1)] + \lambda_0 \mathbb{E}_{z_2 \sim q_2(z_2|x_2)} [\log(1 - D_1(G_1(z_2)))]$$

### Cycle-consistency (CC)

$$\mathcal{L}_{\text{CC}_1}(E_1, G_1, E_2, G_2) = \lambda_3 \text{KL}(q_1(z_1|x_1)||p_\eta(z)) + \lambda_3 \text{KL}(q_2(z_2|x_1^{1 \rightarrow 2}))||p_\eta(z)) - \lambda_4 \mathbb{E}_{z_2 \sim q_2(z_2|x_1^{1 \rightarrow 2})} [\log p_{G_1}(x_1|z_2)]$$

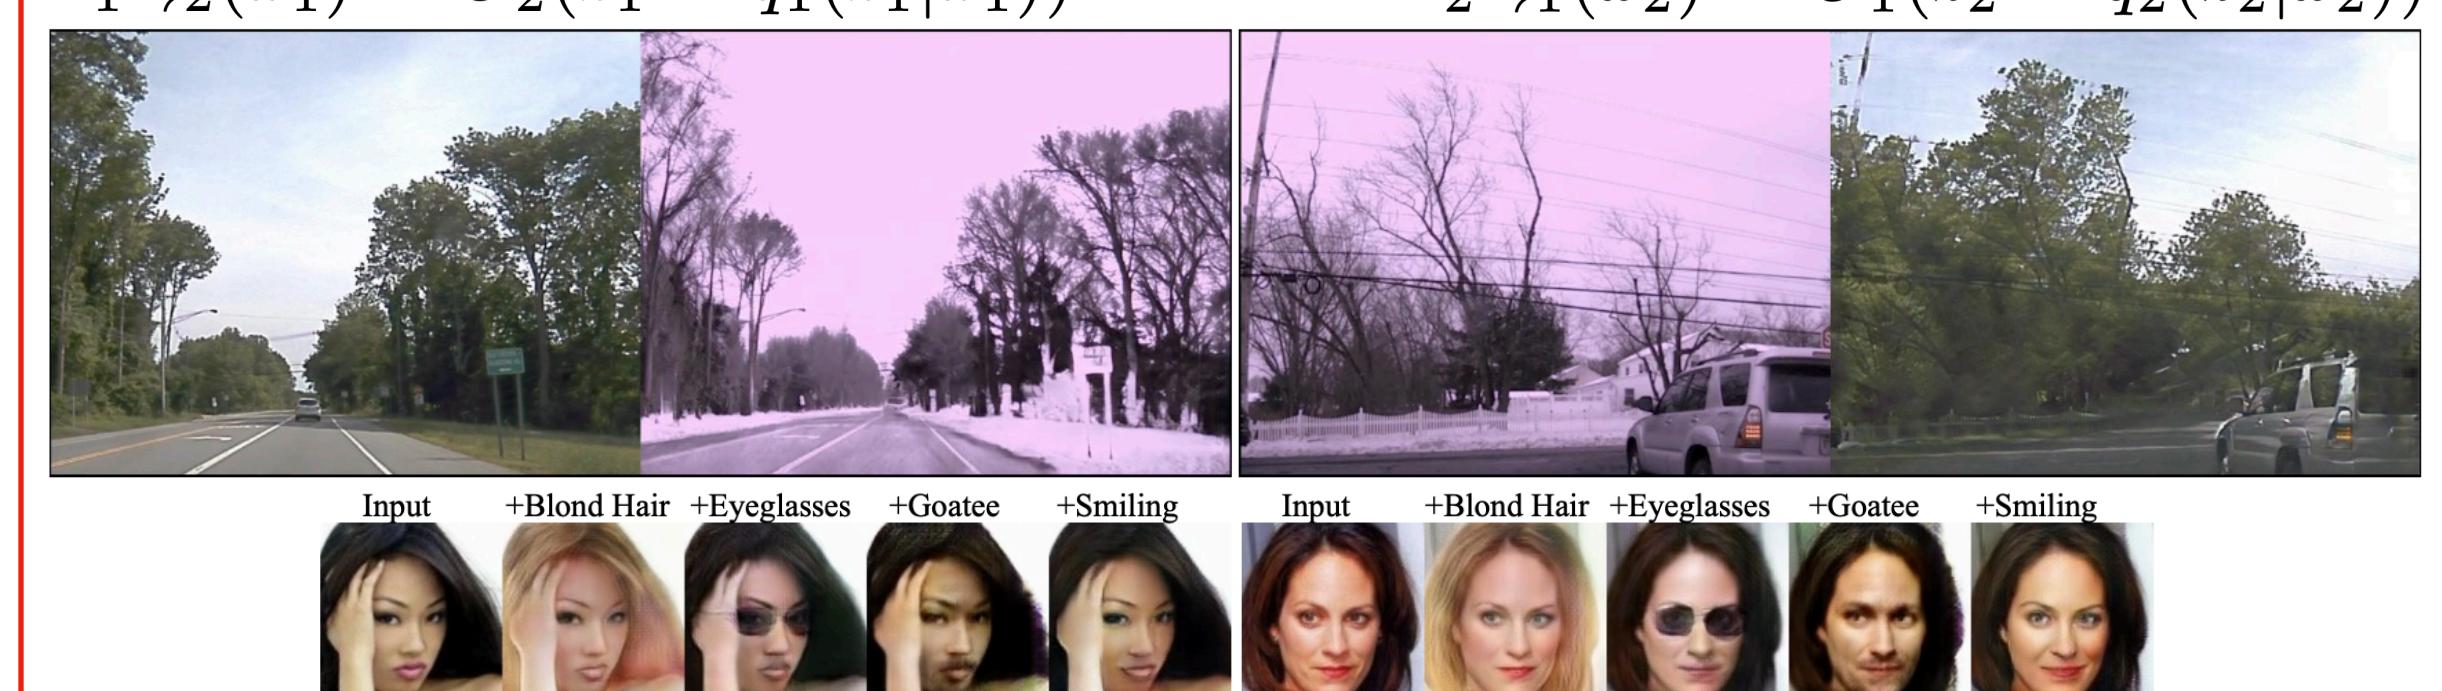
### Learning

$$\min_{E_1, E_2, G_1, G_2} \max_{D_1, D_2} \mathcal{L}_{\text{VAE}_1}(E_1, G_1) + \mathcal{L}_{\text{GAN}_1}(E_1, G_1, D_1) + \mathcal{L}_{\text{CC}_1}(E_1, G_1, E_2, G_2)$$

### Translation

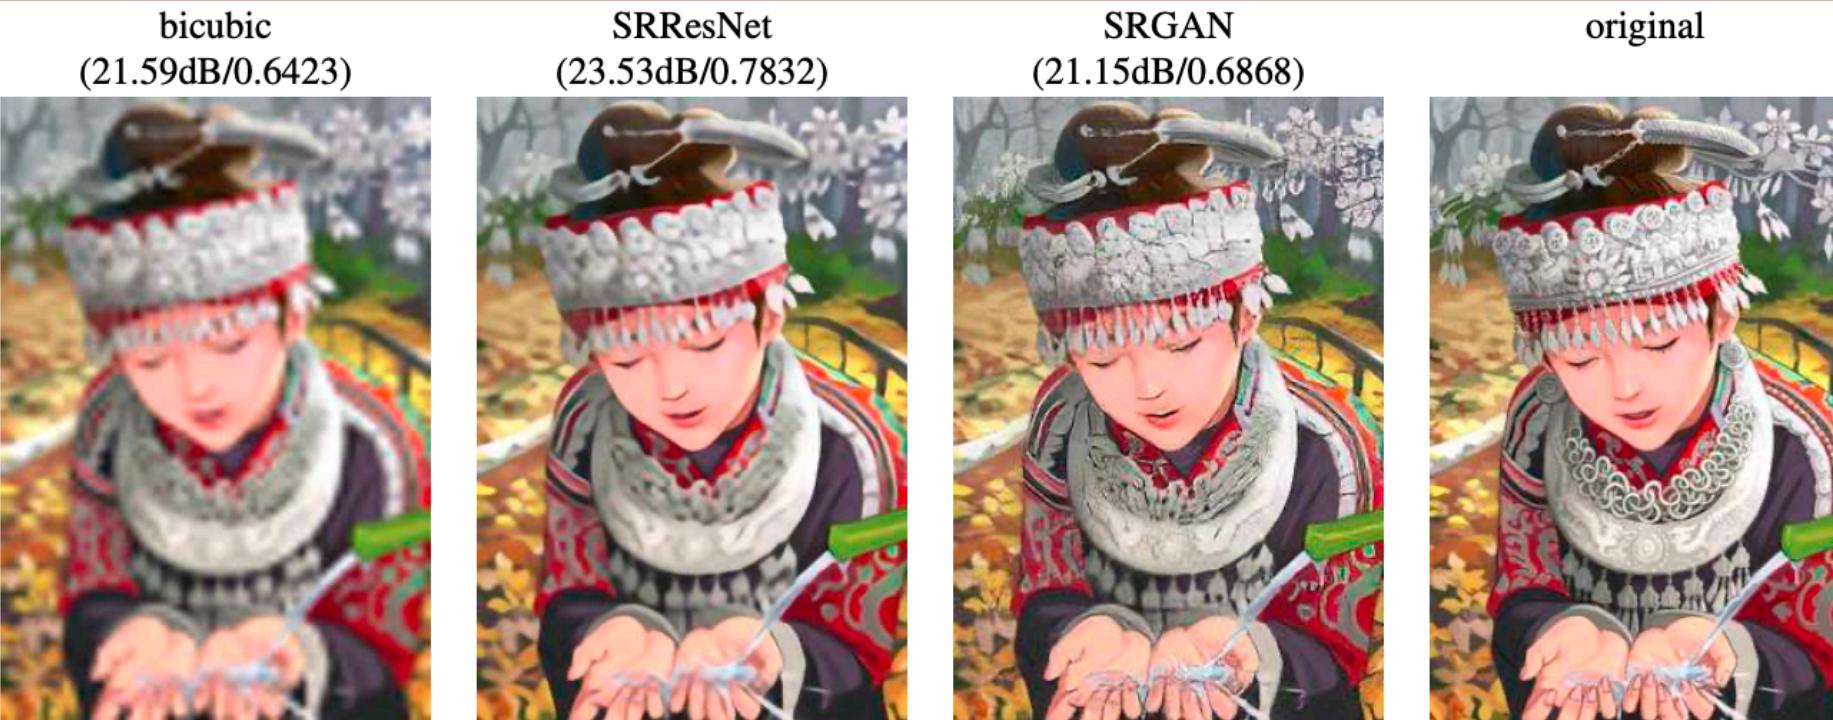
$$F_{1 \rightarrow 2}(x_1) = G_2(z_1 \sim q_1(z_1|x_1))$$

$$F_{2 \rightarrow 1}(x_2) = G_1(z_2 \sim q_2(z_2|x_2))$$





# Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network


[YouTube Video](#)


$I^{SR}$  → high-resolution, super-resolved image

$I^{LR}$  → low-resolution input image

$I^{HR}$  → high-resolution counterpart of  $I^{LR}$

$I^{HR} \triangleright$  Gaussian Filter  $\triangleright$  downsampling (downsampling\_factor =  $r$ )

$C$  → number of color channels

$I^{LR} \in \mathbb{R}^{W \times H \times C}$

$I^{SR}, I^{HR} \in \mathbb{R}^{rW \times rH \times C}$

evaluation metrics →  $\begin{cases} \text{PSNR: peak signal-to-noise ratio} \\ \text{MOS: mean opinion score} \\ \text{SSIM: structural similarity} \end{cases}$

$G_{\theta_G} : I^{LR} \mapsto I^{HR}$  → generator

$\hat{\theta}^G = \arg \min_{\theta_G} \frac{1}{N} \sum_{n=1}^N \ell^{SR}(G_{\theta_G}(I_n^{LR}), I_n^{HR})$  → loss function

$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{\text{train}}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))]$

$$l^{SR} = \underbrace{l_X^{SR}}_{\text{content loss}} + \underbrace{10^{-3} l_{Gen}^{SR}}_{\text{adversarial loss}}$$

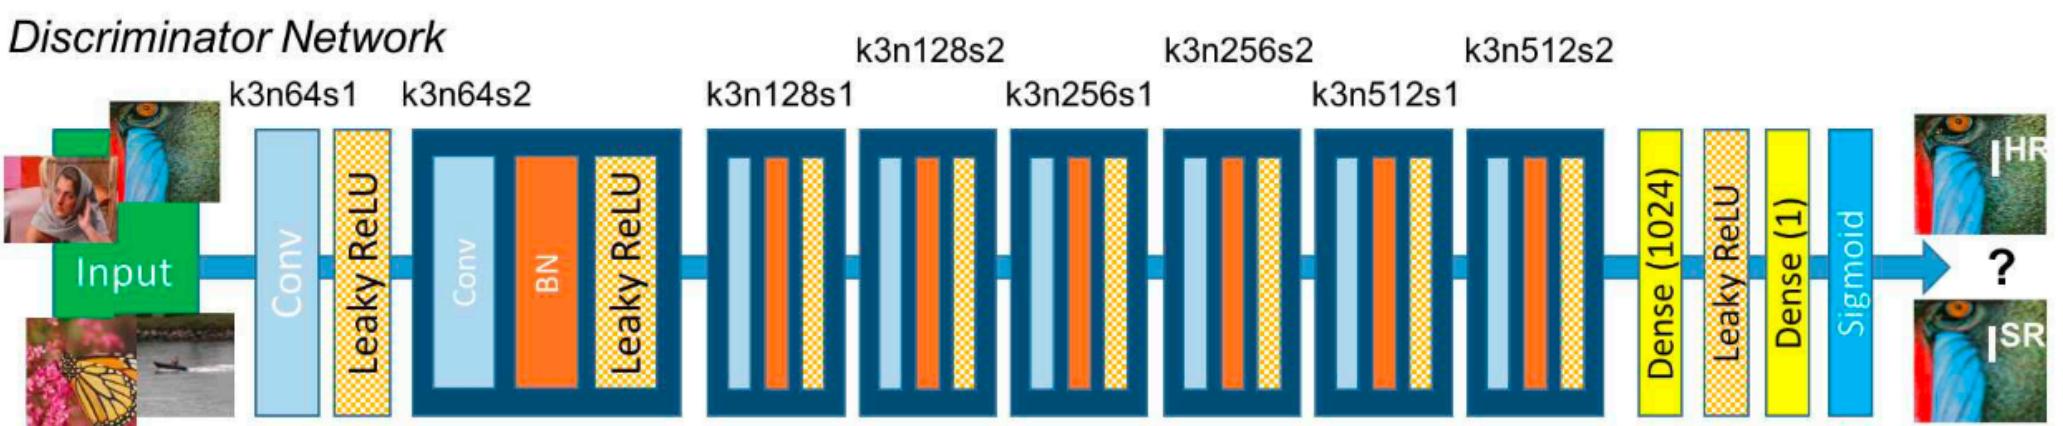
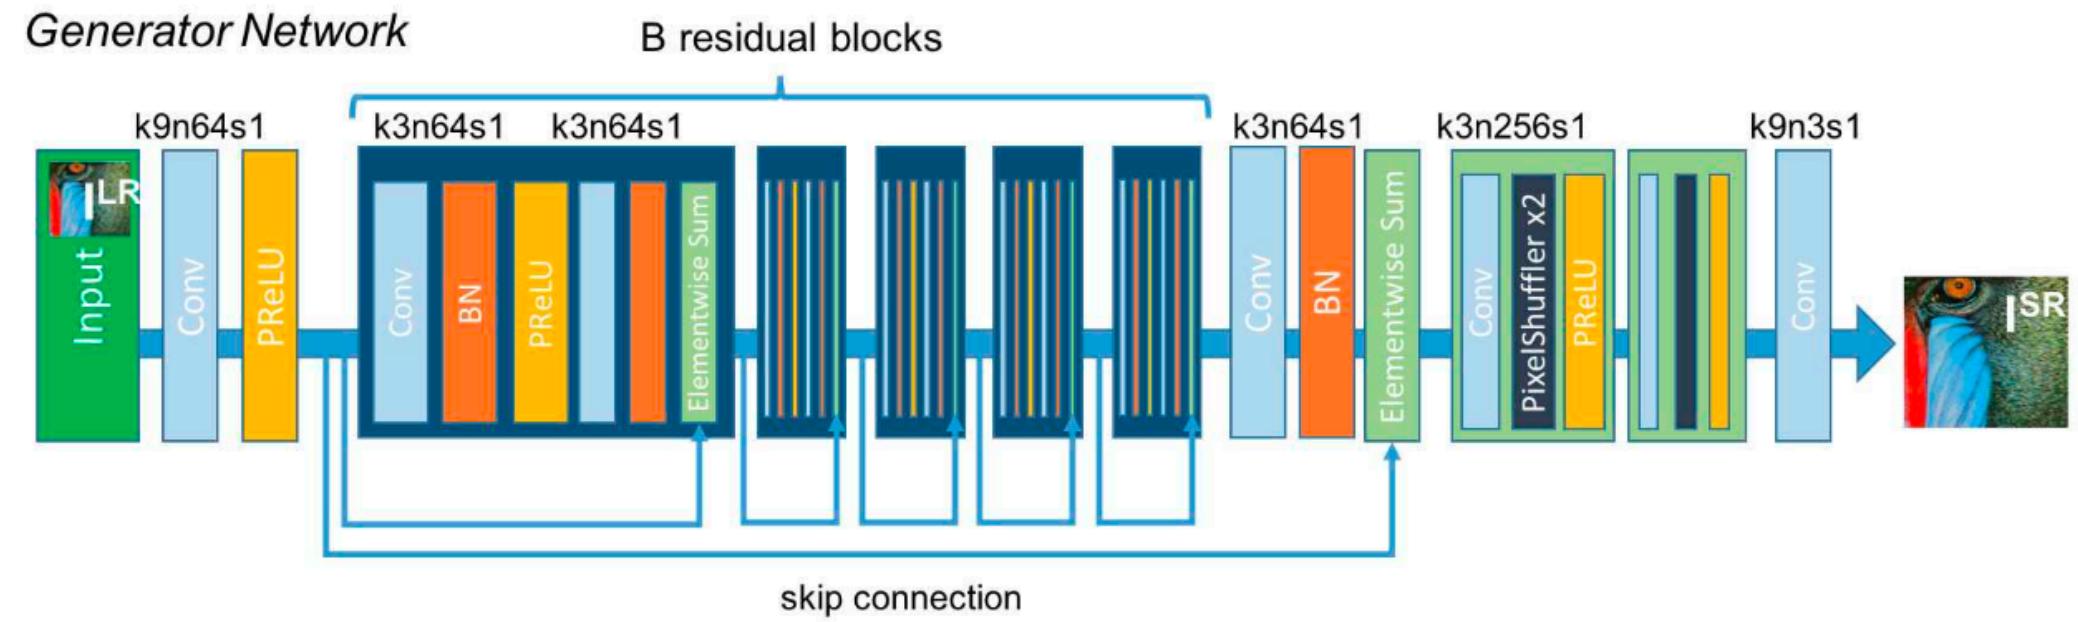
perceptual loss (for VGG based content losses)

$$l_{VGG/i.j}^{SR} = \frac{1}{W_{i,j} H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} (\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}(G_{\theta_G}(I^{LR}))_{x,y})^2$$

$$l_{MSE}^{SR} = \frac{1}{r^2 W H} \sum_{x=1}^{rW} \sum_{y=1}^{rH} (I_{x,y}^{HR} - G_{\theta_G}(I^{LR})_{x,y})^2$$

$$l_{Gen}^{SR} = \sum_{n=1}^N -\log D_{\theta_D}(G_{\theta_G}(I^{LR}))$$

$\phi_{i,j}$  → feature map obtained by the  $j$ -th convolution  
(after activation) before the  $i$ -th maxpooling layer within the VGG19 network





Boulder

# High-Resolution Image Synthesis and Semantic Manipulation with Conditional GANs



[YouTube Video](#)

The pix2pix baseline

$G \rightarrow$  generator

$D \rightarrow$  discriminator

$G \rightarrow$  translate semantic label maps to realistic looking images

$D \rightarrow$  distinguish real images from the translated ones

$\{(s_i, x_i)\} \rightarrow$  training data

$s_i \rightarrow$  semantic label map

$x_i \rightarrow$  corresponding natural photo

$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D)$

$$\mathcal{L}_{\text{GAN}}(G, D) = \mathbb{E}_{(s, x)}[\log D(s, x)] + \mathbb{E}_s[\log(1 - D(s, G(s)))]$$

U-Net as the generator

patch-based fully convolutional network as the discriminator

$256 \times 256$

Improving photorealism and resolution  $2048 \times 1024$

- coarse-to-fine generator
- multi-scale discriminator
- improved adversarial loss

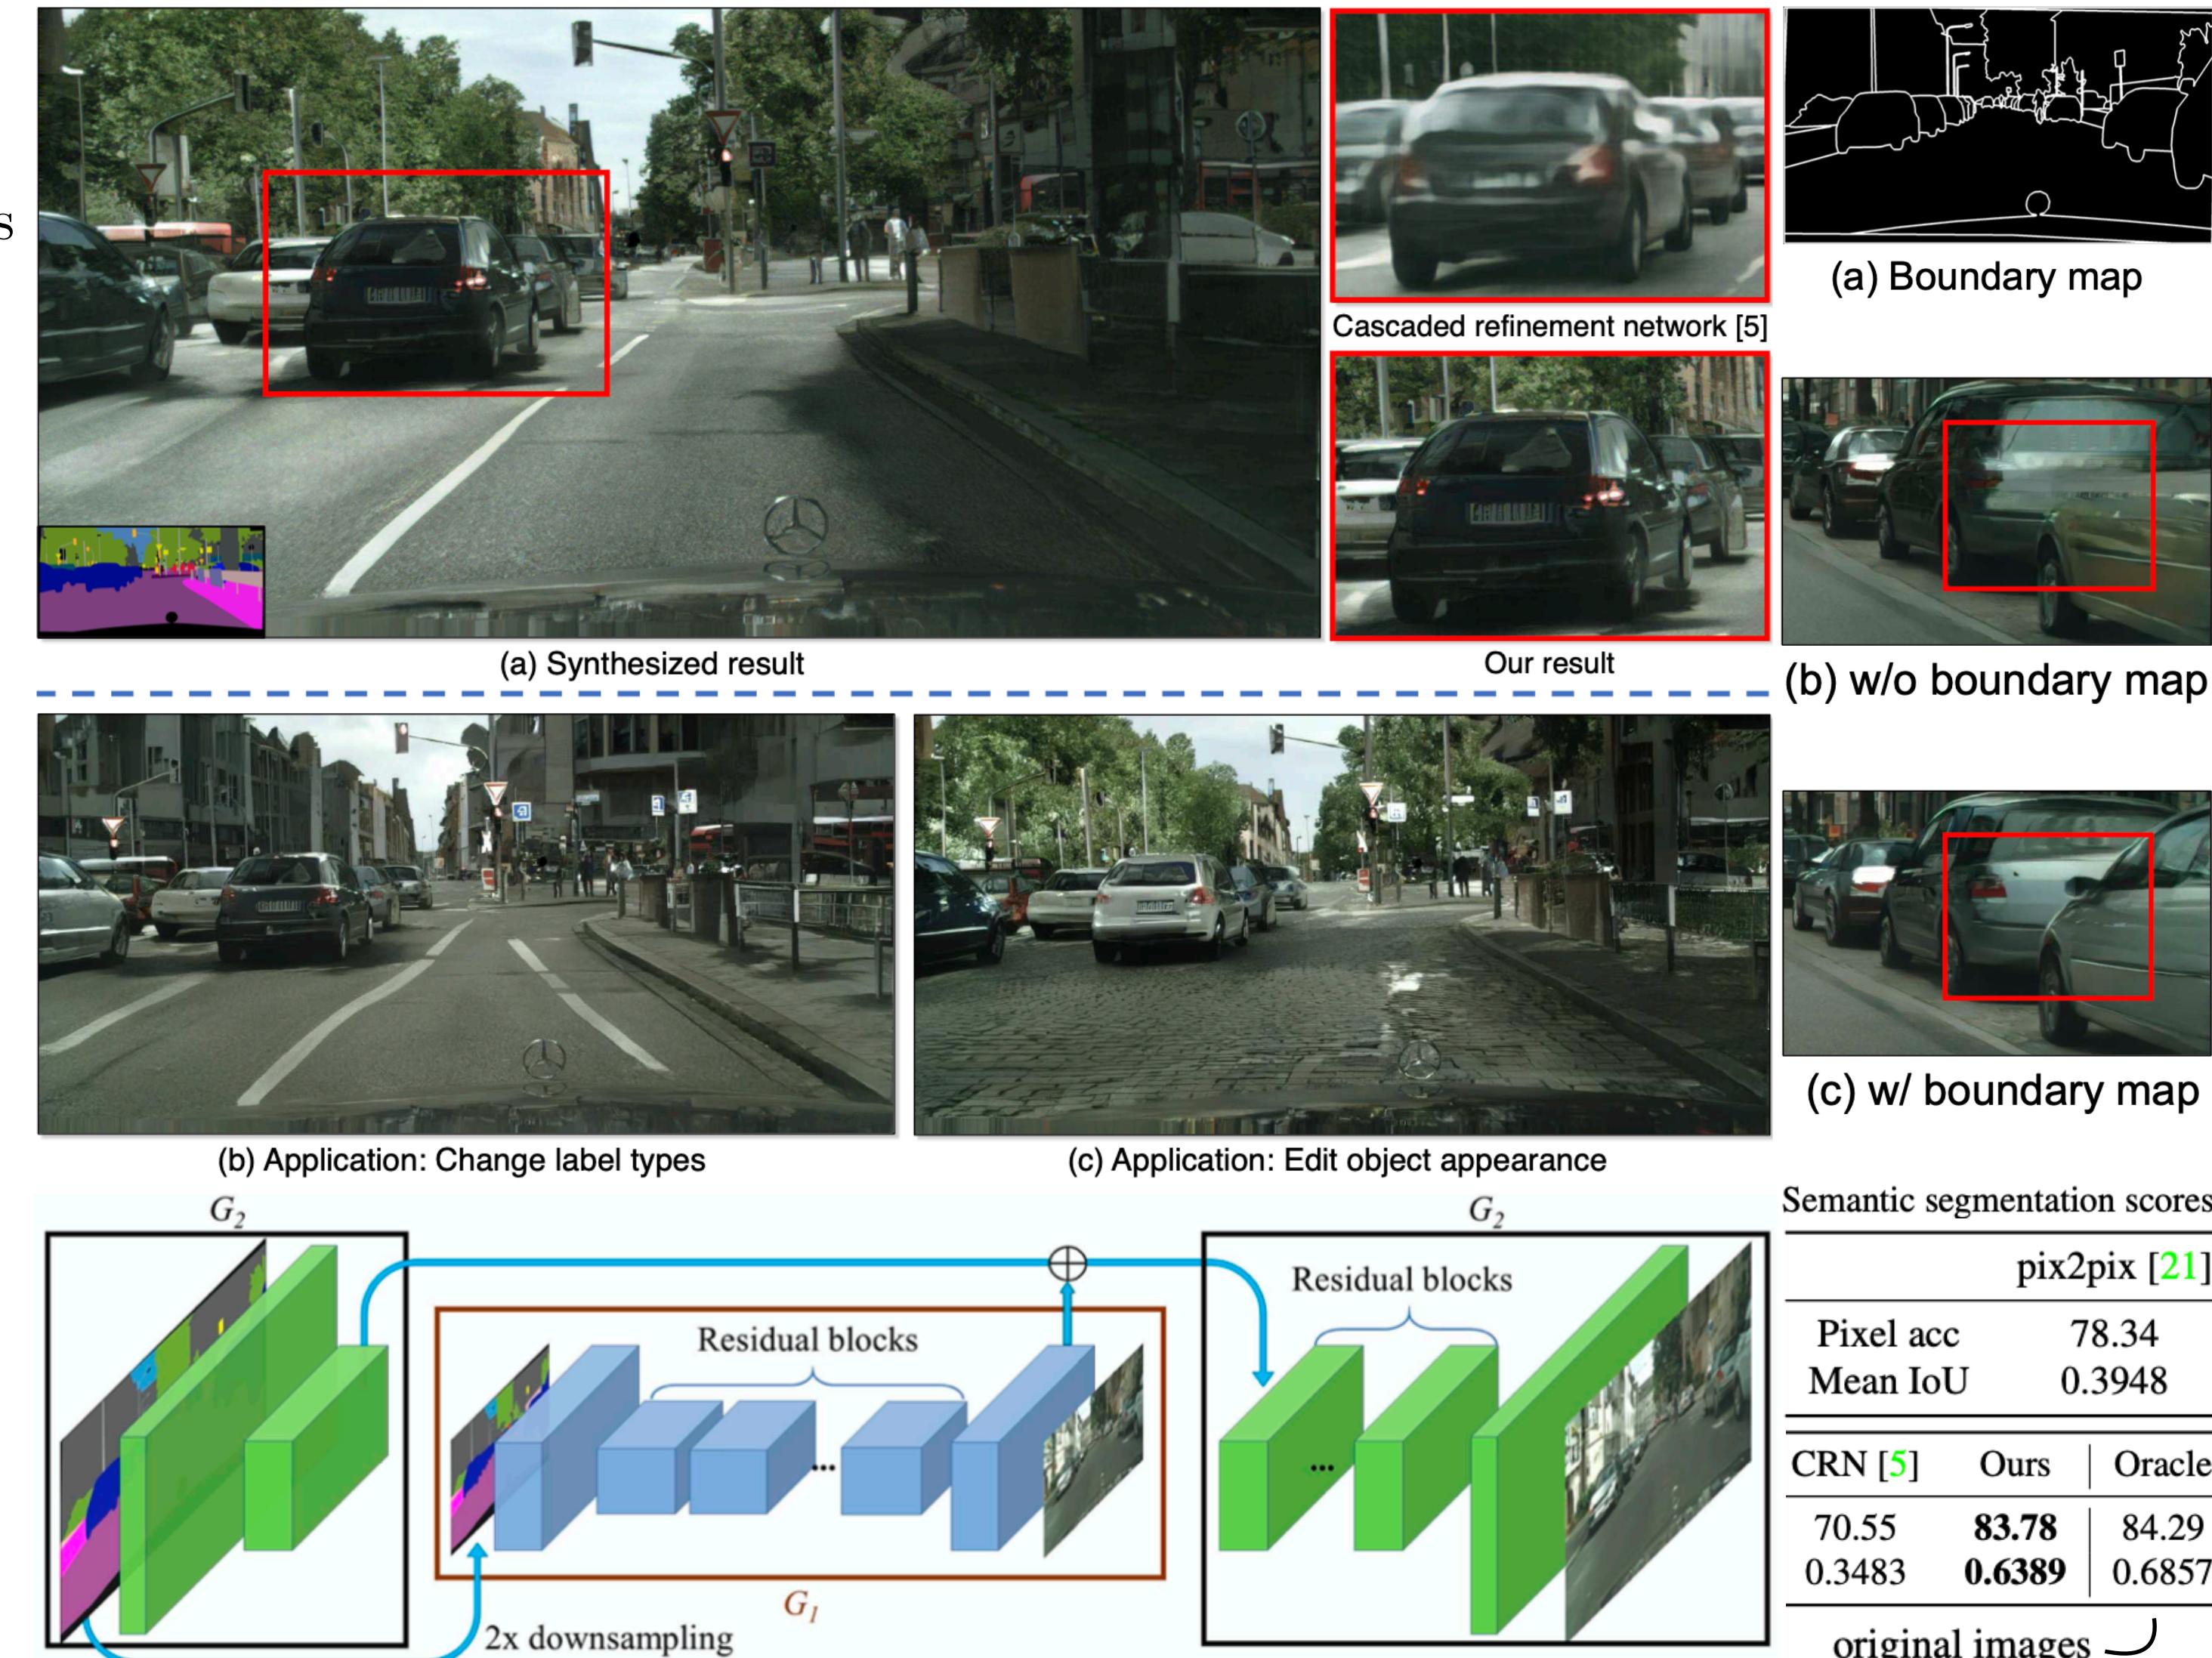
$$\min_G \left( \left( \max_{D_1, D_2, D_3} \sum_{k=1,2,3} \mathcal{L}_{\text{GAN}}(G, D_k) \right) + \lambda \sum_{k=1,2,3} \mathcal{L}_{\text{FM}}(G, D_k) \right)$$

$$\mathcal{L}_{\text{FM}}(G, D_k) = \mathbb{E}_{(\mathbf{s}, \mathbf{x})} \sum_{i=1}^T \frac{1}{N_i} \| | | D_k^{(i)}(\mathbf{s}, \mathbf{x}) - D_k^{(i)}(\mathbf{s}, G(\mathbf{s})) | |_1 \|$$

$D_k^{(i)} \rightarrow$   $i$ -th layer feature extractor of discriminator  $D_k$

$T \rightarrow$  total number of layers

$N_i \rightarrow$  number of elements in each layer

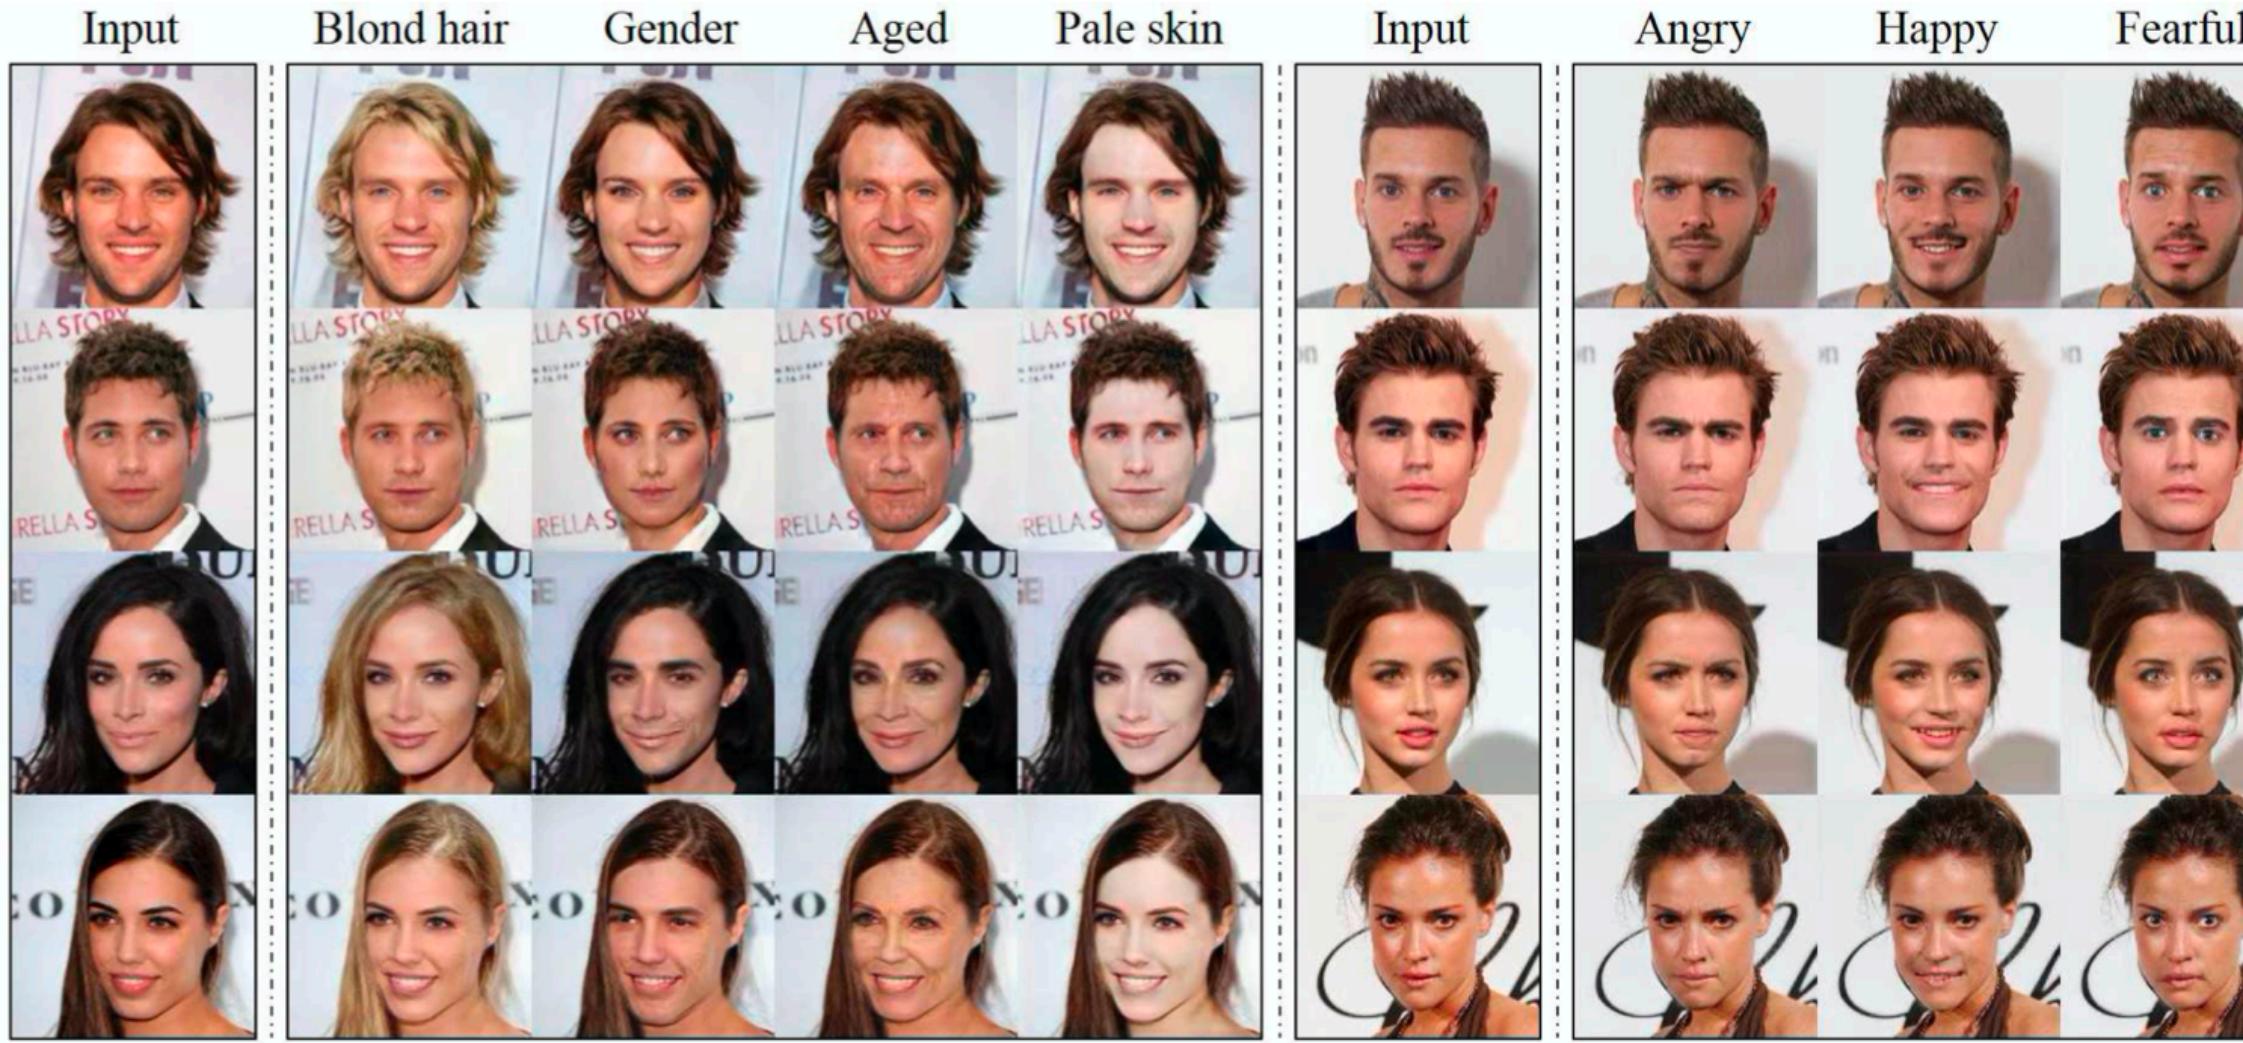




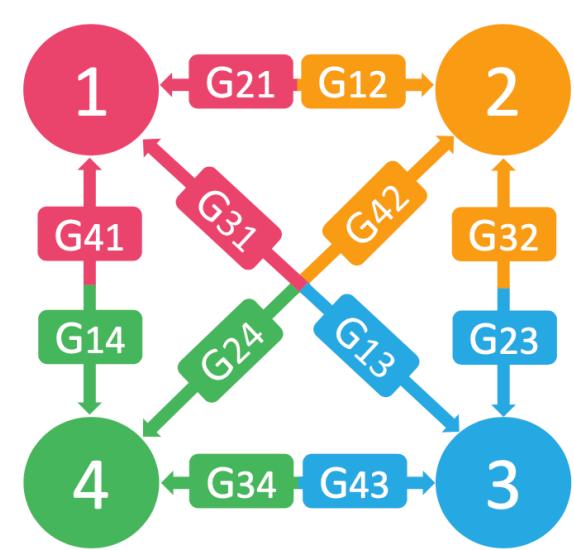
# StarGAN: Unified Generative Adversarial Networks for Multi-Domain Image-to-Image Translation



[YouTube Video](#)

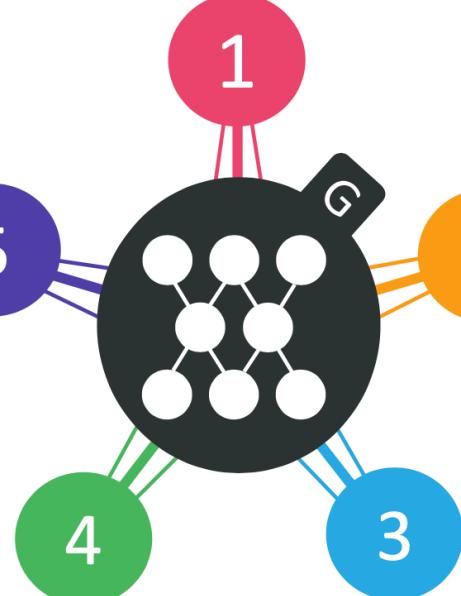


(a) Cross-domain models



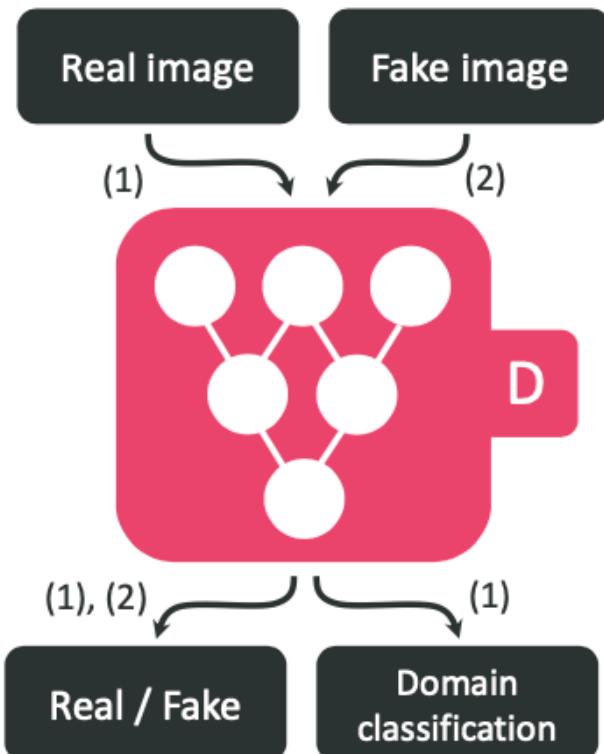
$$\begin{aligned} \mathcal{L}_{adv} &= \mathbb{E}_y[\log D_{src}(y)] + \mathbb{E}_{x,c}[\log(1 - D_{src}(G(x, c)))] \rightarrow \text{adversarial loss} \\ \mathcal{L}_{cls}^r &= \mathbb{E}_{y,c}[-\log D_{cls}(c|y)] \rightarrow \text{domain classification loss of real images} \\ \mathcal{L}_{cls}^f &= \mathbb{E}_{x,c}[-\log D_{cls}(c|G(x, c))] \rightarrow \text{fake images} \\ \mathcal{L}_{rec} &= \mathbb{E}_{x,c,c'}[\|x - G(G(x, c), c')\|_1] \rightarrow \text{reconstruction loss} \\ \lambda_{cls} &= 1 \& \lambda_{rec} = 10 \end{aligned}$$

(b) StarGAN

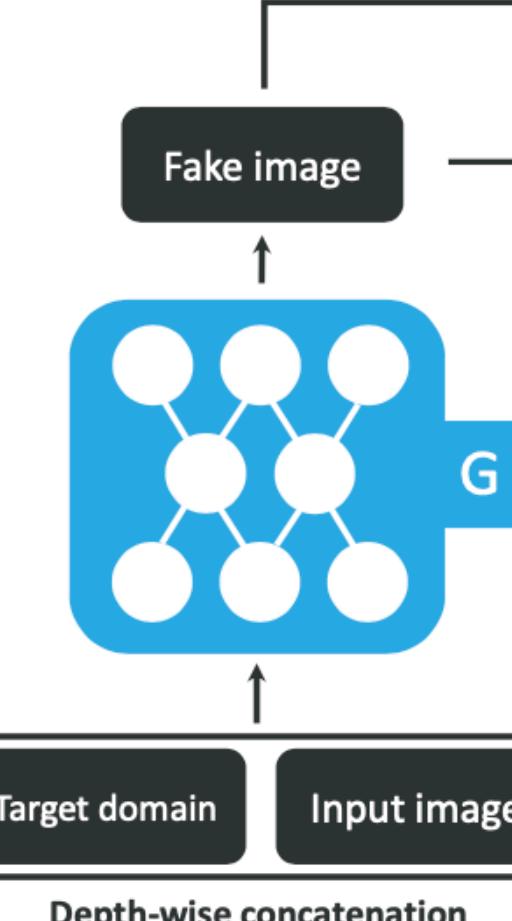


$G \rightarrow$  single generator  
 $x \rightarrow$  input image  
 $y \rightarrow$  output image  
 $c \rightarrow$  target domain label  
 $G(x, c) \rightarrow$  generated image  
 $D : x \mapsto \{D_{src}(x), D_{cls}(x)\}$   
 $\mathcal{L}_D = -\mathcal{L}_{adv} + \lambda_{cls}\mathcal{L}_{cls}^r$   
 $\mathcal{L}_G = \mathcal{L}_{adv} + \lambda_{cls}\mathcal{L}_{cls}^f + \lambda_{rec}\mathcal{L}_{rec}$

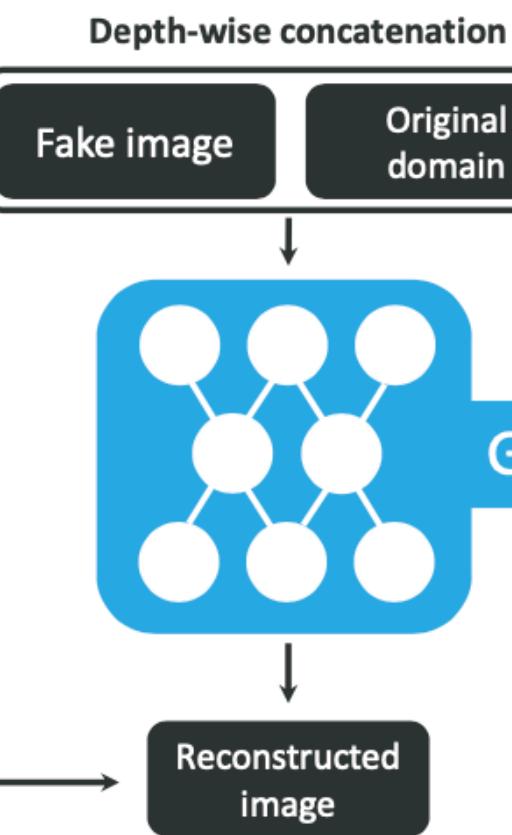
(a) Training the discriminator



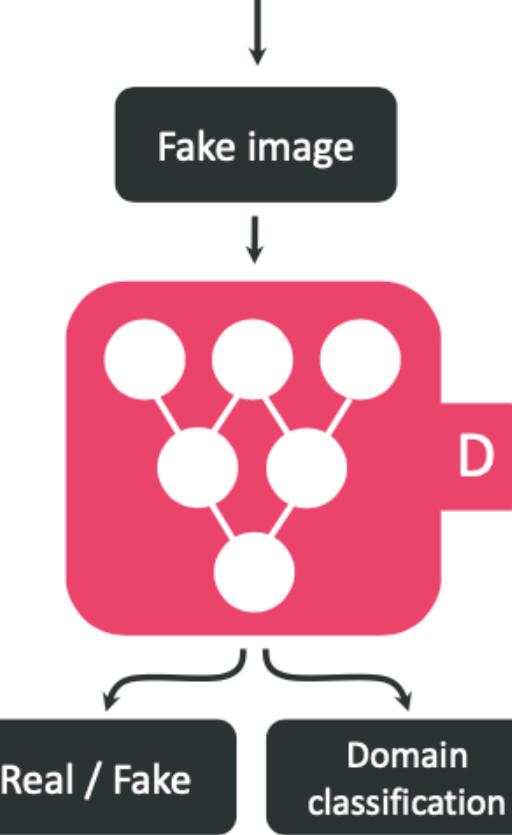
(b) Original-to-target domain



(c) Target-to-original domain



(d) Fooling the discriminator



## Training with multiple datasets

$n \rightarrow$  number of datasets (e.g.,  $n = 2$ )

CelebA and RaFD

$m \rightarrow$  mask vector (n-dimensional one-hot vector)

$\tilde{c} := [c_1, \dots, c_n, m] \rightarrow$  concatenate

$c_i \rightarrow$  vector for labels of the  $i$ -th dataset  
assign zero values to the remaining  $n - 1$  unknown labels

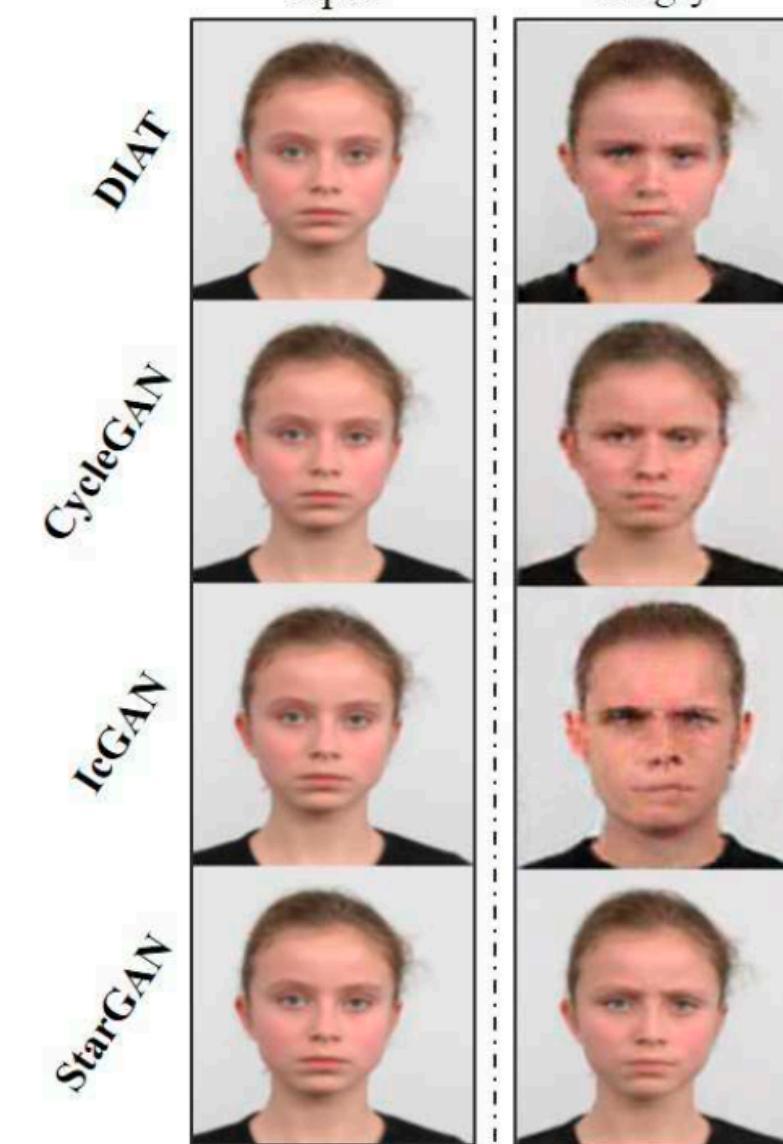
Improved GAN training

$$\mathcal{L}_{adv} = \mathbb{E}_y[D_{src}(y)] - \mathbb{E}_{x,c}[D_{src}(G(x, c))]$$

$$- \lambda_{gp} \mathbb{E}_{\hat{y}}[(\|\nabla_{\hat{y}} D_{src}(\hat{y})\|_2 - 1)^2]$$

$\hat{y} \rightarrow$  sampled uniformly along a straight line between  $y$  and  $G(x, c)$

$$\lambda_{gp} = 10 \rightarrow \text{gradient penalty}$$





Boulder



# Questions?

[YouTube Playlist](#)

---